

Feature and Target Engineering

Data Set

h2o

```
ames <- AmesHousing::make_ames()
ames.h2o <- as.h2o(ames)
```

stratified (*Sale_Price*) training sample

```
set.seed(123)

split <- initial_split(ames, prop = 0.7,
                       strata = "Sale_Price")

ames_train <- training(split)
ames_test <- testing(split)
```

log transformation (*Sale_Price*)

```
ames_recipe <- recipe(Sale_Price ~ ., data = ames_train) %>%
  step_log(all_outcomes())

ames_recipe
```

Data Recipe

Inputs:

	role	#variables
outcome		1
predictor		80

Operations:

Log transformation on *all_outcomes*

Box-Cox transformation (example)

```
lambda <- 3

y <- forecast::BoxCox(10, lambda)

inv_box_cox <- function(x, lambda) {
  # for Box-Cox, lambda = 0 -> log transform
  if(lambda == 0) exp(x) else (lambda*x + 1)^(1/lambda)
```

```

}

inv_box_cox(y, lambda)

[1] 10
attr(,"lambda")
[1] 3

# Log transformation
train_log_y <- log(ames_train$Sale_Price)
test_log_y  <- log(ames_train$Sale_Price)

# Box Cox transformation
lambda <- forecast::BoxCox.lambda(ames_train$Sale_Price)
train_bc_y <- forecast::BoxCox(ames_train$Sale_Price, lambda)
test_bc_y  <- forecast::BoxCox(ames_test$Sale_Price, lambda)

# Plot differences
levs <- c("Normal", "Log_Transform", "BoxCox_Transform")
data.frame(
  Normal = ames_train$Sale_Price,
  Log_Transform = train_log_y,
  BoxCox_Transform = train_bc_y
) %>%
  gather(Transform, Value) %>%
  mutate(Transform = factor(Transform, levels = levs)) %>%
  ggplot(aes(Value, fill = Transform)) +
    geom_histogram(show.legend = FALSE, bins = 40) +
    facet_wrap(~ Transform, scales = "free_x")

```





Figure 1: Response variable transformations.

Missing Values

```
sum(is.na(AmesHousing::ames_raw))
```

```
[1] 13997
```

```
AmesHousing::ames_raw %>%
  is.na() %>%
  reshape2::melt() %>%
  ggplot(aes(Var2, Var1, fill = value)) +
    geom_raster() +
    coord_flip() +
    scale_y_continuous(NULL, expand = c(0,0)) +
    scale_fill_grey(name = "",
                    labels = c("Present",
                              "Missing")) +
  xlab("Observations") +
  theme(axis.text.y = element_text(size = 4))
```



Missing Garage?

```
AmesHousing::ames_raw %>%
  filter(is.na(`Garage Type`)) %>%
  select(starts_with("Garage"))
```

```
# A tibble: 157 x 7
```

```
  `Garage Type` `Garage Yr Blt` `Garage Finish` `Garage Cars` `Garage Area`
  <chr>          <int> <chr>          <int>          <int>
```


```

1 <NA>      NA <NA>      0      0
2 <NA>      NA <NA>      0      0
3 <NA>      NA <NA>      0      0
4 <NA>      NA <NA>      0      0
5 <NA>      NA <NA>      0      0
6 <NA>      NA <NA>      0      0
7 <NA>      NA <NA>      0      0
8 <NA>      NA <NA>      0      0
9 <NA>      NA <NA>      0      0
10 <NA>     NA <NA>      0      0
# ... with 147 more rows, and 2 more variables: `Garage Qual` <chr>, `Garage
#   Cond` <chr>

```

Missing values w/cluster (*visdat*)

```
vis_miss(AmesHousing::ames_raw, cluster = T)
```



Missing Value Imputation

basic descriptive statistic

```
ames_recipe %>%
  step_medianimpute(Gr_Liv_Area)
```

Data Recipe

Inputs:

```

    role #variables
outcome          1
predictor        80

```

Operations:

Log transformation on all_outcomes

Median Imputation for Gr_Liv_Area

KNN approach (typical k = 5-10)

```

ames_recipe %>%
  step_knnimpute(all_predictors(), neighbors = 6)

```

Data Recipe

Inputs:

```

    role #variables
outcome          1
predictor        80

```

Operations:

Log transformation on all_outcomes

K-nearest neighbor imputation for all_predictors

```

impute_ames <- ames_train
set.seed(123)
index <- sample(seq_along(impute_ames$Gr_Liv_Area), 50)
actuals <- ames_train[index, ]
impute_ames$Gr_Liv_Area[index] <- NA

p1 <- ggplot() +
  geom_point(data = impute_ames, aes(Gr_Liv_Area, Sale_Price), alpha = .2) +
  geom_point(data = actuals, aes(Gr_Liv_Area, Sale_Price), color = "red") +
  scale_x_log10(limits = c(300, 5000)) +
  scale_y_log10(limits = c(10000, 500000)) +
  ggtitle("Actual values")

# Mean imputation
mean_juiced <- recipe(Sale_Price ~ ., data = impute_ames) %>%
  step_meanimpute(Gr_Liv_Area) %>%
  prep(training = impute_ames, retain = TRUE) %>%

```

```
juice()
mean_impute <- mean_juiced[index, ]

p2 <- ggplot() +
  geom_point(data = actuals, aes(Gr_Liv_Area, Sale_Price), color = "red") +
  geom_point(data = mean_impute, aes(Gr_Liv_Area, Sale_Price), color = "blue") +
  scale_x_log10(limits = c(300, 5000)) +
  scale_y_log10(limits = c(10000, 500000)) +
  ggtitle("Mean Imputation")

# KNN imputation
knn_juiced <- recipe(Sale_Price ~ ., data = impute_ames) %>%
  step_knnimpute(Gr_Liv_Area) %>%
  prep(training = impute_ames, retain = TRUE) %>%
  juice()
knn_impute <- knn_juiced[index, ]

p3 <- ggplot() +
  geom_point(data = actuals, aes(Gr_Liv_Area, Sale_Price), color = "red") +
  geom_point(data = knn_impute, aes(Gr_Liv_Area, Sale_Price), color = "blue") +
  scale_x_log10(limits = c(300, 5000)) +
  scale_y_log10(limits = c(10000, 500000)) +
  ggtitle("KNN Imputation")

# Bagged imputation
bagged_juiced <- recipe(Sale_Price ~ ., data = impute_ames) %>%
  step_bagimpute(Gr_Liv_Area) %>%
  prep(training = impute_ames, retain = TRUE) %>%
  juice()
bagged_impute <- bagged_juiced[index, ]

p4 <- ggplot() +
  geom_point(data = actuals, aes(Gr_Liv_Area, Sale_Price), color = "red") +
  geom_point(data = bagged_impute, aes(Gr_Liv_Area, Sale_Price), color = "blue") +
  scale_x_log10(limits = c(300, 5000)) +
  scale_y_log10(limits = c(10000, 500000)) +
  ggtitle("Bagged Trees Imputation")

gridExtra::grid.arrange(p1, p2, p3, p4, nrow = 2)
```

Warning: Removed 63 rows containing missing values (geom_point).

Increase in training time by model type:

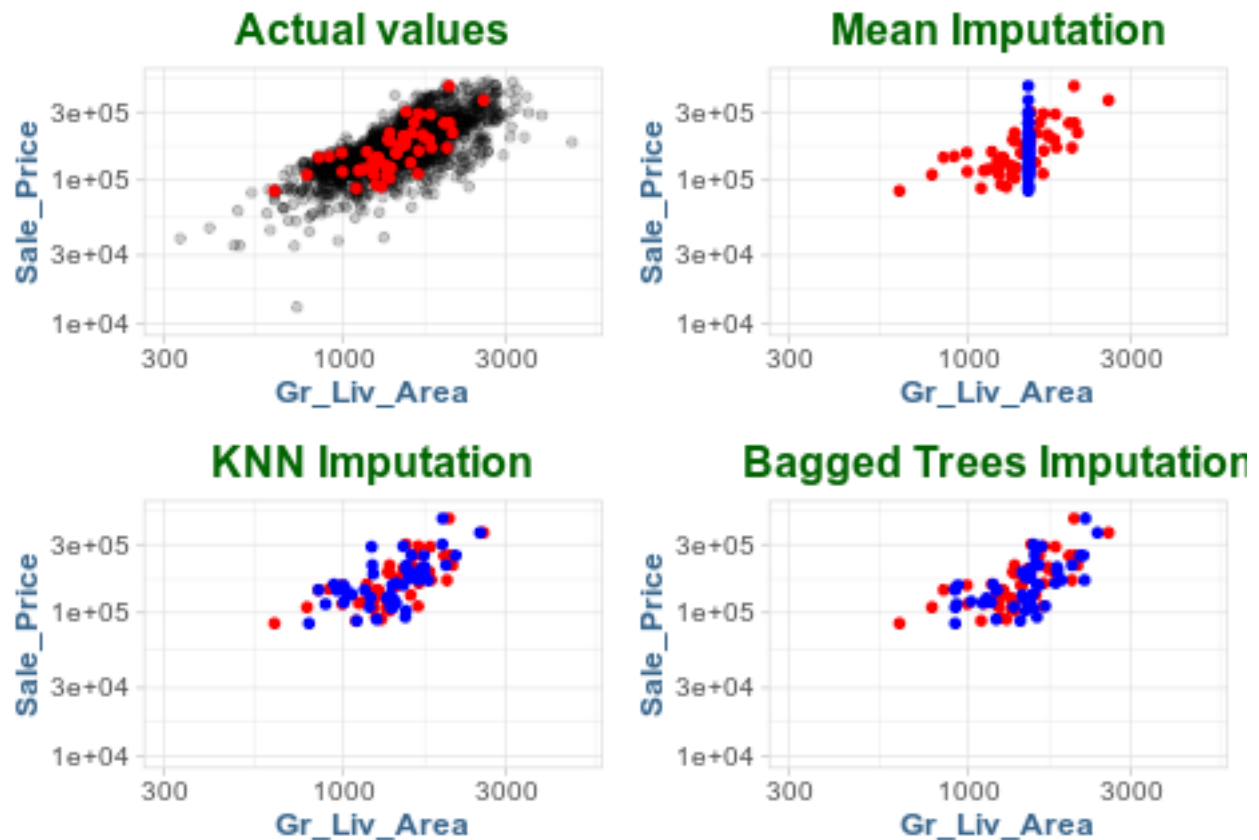


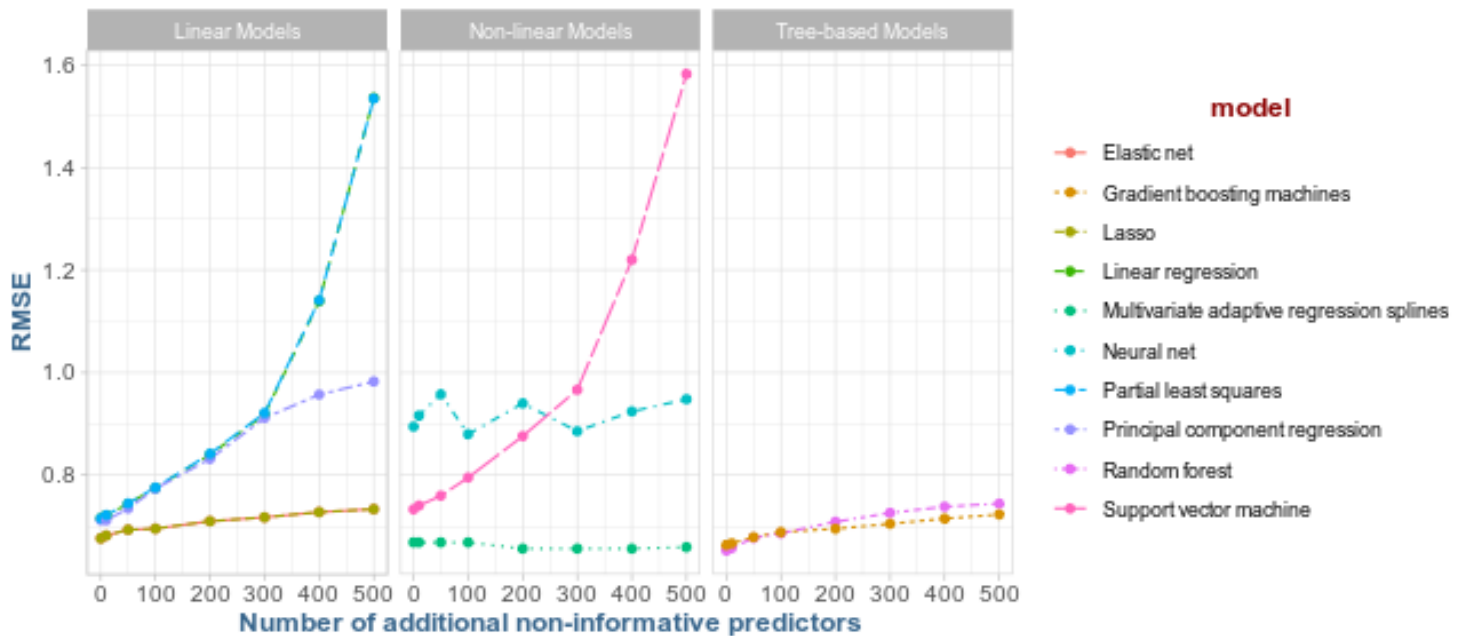
Figure 2: Comparison of three different imputation methods. The red points represent actual values which were removed and made missing and the blue points represent the imputed values. Estimated statistic imputation methods (i.e. mean, median) merely predict the same value for each observation and can reduce the signal between a feature and the response; whereas KNN and tree-based procedures tend to maintain the feature distribution and relationship.

```
model_results <- read_csv(paste0(data.dir, "feature-selection-impacts-results.csv")) %>%
  mutate(type = case_when(
    model %in% c("lm", "pcr", "pls", "glmnet", "lasso") ~ "Linear Models",
    model %in% c("earth", "svmLinear", "nn") ~ "Non-linear Models",
    TRUE ~ "Tree-based Models"
  )) %>%
  mutate(model = case_when(
    model == "lm" ~ "Linear regression",
    model == "earth" ~ "Multivariate adaptive regression splines",
    model == "gbm" ~ "Gradient boosting machines",
    model == "glmnet" ~ "Elastic net",
    model == "lasso" ~ "Lasso",
    model == "nn" ~ "Neural net",
    model == "pcr" ~ "Principal component regression",
    model == "pls" ~ "Partial least squares",
    model == "ranger" ~ "Random forest",
    TRUE ~ "Support vector machine"
  ))
```

Parsed with column specification:

```
cols(
  model = col_character(),
  NIP = col_double(),
  RMSE = col_double(),
  time = col_double()
)

ggplot(model_results, aes(NIP, RMSE, color = model, lty = model)) +
  geom_line() +
  geom_point() +
  facet_wrap(~ type, nrow = 1) +
  xlab("Number of additional non-informative predictors")
```

```
model_results %>%
  group_by(model) %>%
  mutate(
    time_impact = time / first(time),
    time_impact = time_impact - 1
  ) %>%
  ggplot(aes(NIP, time_impact, color = model, lty = model)) +
    geom_line() +
    geom_point() +
    facet_wrap(~ type, nrow = 1) +
    scale_y_continuous("Percent increase in training time",
                      labels = scales::percent) +
    xlab("Number of additional non-informative predictors")
```

Rules of thumb for zero variance features:

- The fraction of unique values over the sample size is low (say < 10%)
- The ratio of the frequency of the most prevalent value to the frequency of the second most prevalent value is large (say > 20%)

If both of these criteria are met, then it is often advantageous to remove them from the model.

```
caret::nearZeroVar(ames_train, saveMetrics = T) %>%
  rownames_to_column() %>%
  filter(nzv)
```

	rowname	freqRatio	percentUnique	zeroVar	nzv
1	Street	292.28571	0.09741841	FALSE	TRUE
2	Alley	20.52688	0.14612762	FALSE	TRUE

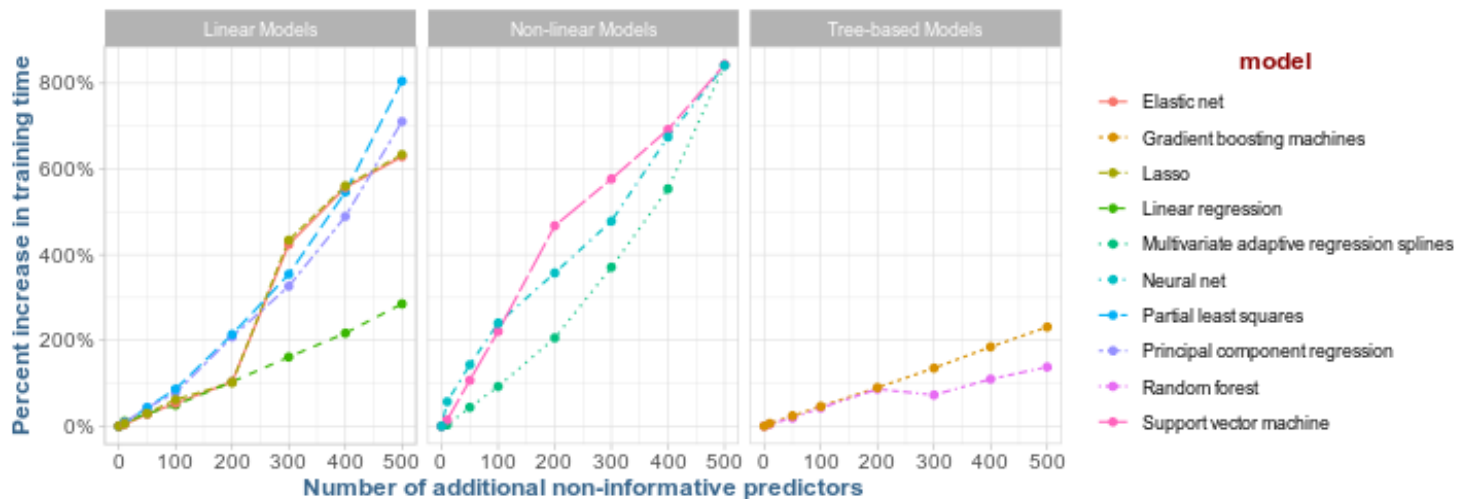


Figure 3: Impact in model training time as non-informative predictors are added.

3	Land_Contour	22.28916	0.19483682	FALSE	TRUE
4	Utilities	1025.00000	0.14612762	FALSE	TRUE
5	Land_Slope	22.76744	0.14612762	FALSE	TRUE
6	Condition_2	203.10000	0.34096444	FALSE	TRUE
7	Roof_Mat1	126.50000	0.24354603	FALSE	TRUE
8	Bsmt_Cond	19.93478	0.29225524	FALSE	TRUE
9	BsmtFin_Type_2	21.50617	0.34096444	FALSE	TRUE
10	Heating	101.05000	0.24354603	FALSE	TRUE
11	Low_Qual_Fin_SF	1013.00000	1.31514856	FALSE	TRUE
12	Kitchen_AbvGr	23.68675	0.19483682	FALSE	TRUE
13	Functional	38.18000	0.34096444	FALSE	TRUE
14	Enclosed_Porch	100.94118	7.40379932	FALSE	TRUE
15	Three_season_porch	674.66667	1.16902094	FALSE	TRUE
16	Screen_Porch	234.87500	4.52995616	FALSE	TRUE
17	Pool_Area	2045.00000	0.43838285	FALSE	TRUE
18	Pool_QC	681.66667	0.24354603	FALSE	TRUE
19	Misc_Feature	30.49231	0.19483682	FALSE	TRUE
20	Misc_Val	165.33333	1.41256698	FALSE	TRUE

Numeric Feature Engineering

Skewness can have a drastic impact on the performance of GLMs & regularized models.

Non-parametric models are rarely affected by skewed features; however, normalizing features will not have a negative effect on these models' performance. For example, normalizing features will only shift the optimal split points in tree-based algorithms. Consequently, when in doubt, normalize.

Skewness

Data Recipe

Inputs:

role	#variables
outcome	1
predictor	80

Operations:

Yeo-Johnson transformation on all_numeric

Standardization

```
ames_recipe %>%  
  step_center(all_numeric(), -all_outcomes()) %>%  
  step_scale(all_numeric(), -all_outcomes())
```

Data Recipe

Inputs:

role	#variables
outcome	1
predictor	80

Operations:

Log transformation on all_outcomes

Centering for all_numeric, -, all_outcomes()

Scaling for all_numeric, -, all_outcomes()

```
set.seed(123)  
x1 <- tibble(  
  variable = "x1",  
  `Real Value` = runif(25, min = -30, max = 5),  
  `Standardized Value` = scale(`Real Value`) %>% as.numeric()  
)  
  
set.seed(456)  
x2 <- tibble(  
  variable = "x1",  
  `Real Value` = runif(25, min = -30, max = 5),  
  `Standardized Value` = scale(`Real Value`) %>% as.numeric()  
)
```

```

variable = "x2",
`Real value` = rlnorm(25, log(25)),
`Standardized value` = scale(`Real value`) %>% as.numeric()
)

set.seed(789)
x3 <- tibble(
  variable = "x3",
  `Real value` = rnorm(25, 150, 15),
  `Standardized value` = scale(`Real value`) %>% as.numeric()
)

x1 %>%
  bind_rows(x2) %>%
  bind_rows(x3) %>%
  gather(key, value, -variable) %>%
  mutate(variable = factor(variable, levels = c("x3", "x2", "x1"))) %>%
  ggplot(aes(value, variable)) +
    geom_point(alpha = .6) +
    facet_wrap(~ key, scales = "free_x") +
    ylab("Feature") +
    xlab("Value")

```

Warning: Removed 150 rows containing missing values (geom_point).




Figure 4: Standardizing features allows all features to be compared on a common value scale regardless of their real value differences.

Categorical Feature Engineering

Lumping

When a feature contains levels that have few observations.

For example:

```
count(ames_train, Neighborhood) %>% arrange(n)
```

```
# A tibble: 27 x 2
  Neighborhood      n
  <fct>           <int>
1 Green_Hills      2
2 Greens           7
3 Blueste          8
4 Northpark_Villa 17
5 Briardale        18
6 Veenker          20
7 Bloomington_Heights 21
8 South_and_West_of_Iowa_State_University 27
9 Meadow_Village   29
10 Clear_Creek     31
# ... with 17 more rows
```

```
count(ames_train, Screen_Porch) %>% arrange(n)
```

```
# A tibble: 93 x 2
  Screen_Porch      n
  <int> <int>
1      40      1
2      63      1
3      80      1
4      92      1
5      94      1
6      99      1
7     104      1
8     109      1
9     110      1
10     111      1
# ... with 83 more rows
```

We can benefit from lumping these together into an “other” category when they contain less than 10% of the training sample.

Note: This can have an adverse effect on performance

```
lumping <- recipe(Sale_Price ~., data = ames_train) %>%
  step_other(Neighborhood, threshold = 0.01,
             other = "other") %>%
  step_other(Screen_Porch, threshold = 0.1,
             other = ">0")

apply_2_training <- prep(lumping, training = ames_train) %>%
  bake(ames_train)

# New distribution of Neighborhood
count(apply_2_training, Neighborhood) %>% arrange(n)
```

```
# A tibble: 22 x 2
  Neighborhood      n
  <fct>           <int>
1 Bloomington_Heights 21
2 South_and_West_of_Iowa_State_University 27
3 Meadow_Village      29
4 Clear_Creek         31
5 Stone_Brook         34
6 Northridge         48
7 Timberland         55
8 Iowa_DOT_and_Rail_Road 62
9 Crawford           72
10 other             72
# ... with 12 more rows
```

```
# New distribution of Screen_Porch
count(apply_2_training, Screen_Porch) %>% arrange(n)
```

```
# A tibble: 2 x 2
  Screen_Porch      n
  <fct>         <int>
1 >0           174
2 0           1879
```

```
dat <- data.table(id = 1:9, x = rep(c("a", "b", "c"), 3))
dat
```

```
id x
1: 1 a
2: 2 b
3: 3 c
4: 4 a
5: 5 b
6: 6 c
```

```
7: 7 a
8: 8 b
9: 9 c
```

```
# full-rank
```

```
dat[, .(id,
  `X = a` = as.numeric(x == "a"),
  `X = b` = as.numeric(x == "b"),
  `X = c` = as.numeric(x == "c")) ]
```

```
   id X = a X = b X = c
1:  1     1     0     0
2:  2     0     1     0
3:  3     0     0     1
4:  4     1     0     0
5:  5     0     1     0
6:  6     0     0     1
7:  7     1     0     0
8:  8     0     1     0
9:  9     0     0     1
```

```
# one-hot (leave one out)
```

```
dat[, .(id,
  `X = a` = as.numeric(x == "a"),
  `X = b` = as.numeric(x == "b")) ]
```

```
   id X = a X = b
1:  1     1     0
2:  2     0     1
3:  3     0     0
4:  4     1     0
5:  5     0     1
6:  6     0     0
7:  7     1     0
8:  8     0     1
9:  9     0     0
```

```
# Lump levels for two features
```

```
recipe(Sale_Price ~., data = ames_train) %>%
  step_dummy(all_nominal(), one_hot = T)
```

Data Recipe

Inputs:

```
   role #variables
outcome          1
```

predictor 80

Operations:

Dummy variables from `all_nominal`

Label Encoding

Label encoding is a pure numeric conversion of the levels of a categorical variable.

For example, the `MS_SubClass` variable has 16 levels, which we can reencode numerically.

Important: The features will be treated as ordered (ordinal encoding), so if the feature is not naturally ordered, this will have a poor impact on the model.

```
count(ames_train, MS_SubClass)
```

```
# A tibble: 16 x 2
  MS_SubClass      n
  <fct>          <int>
1 One_Story_1946_and_Newer_All_Styles    753
2 One_Story_1945_and_Older                91
3 One_Story_with_Finished_Attic_All_Ages    5
4 One_and_Half_Story_Unfinished_All_Ages   11
5 One_and_Half_Story_Finished_All_Ages   211
6 Two_Story_1946_and_Newer             395
7 Two_Story_1945_and_Older              98
8 Two_and_Half_Story_All_Ages            17
9 Split_or_Multilevel                   75
10 Split_Foyer                          32
11 Duplex_All_Styles_and_Ages             66
12 One_Story_PUD_1946_and_Newer          145
13 One_and_Half_Story_PUD_All_Ages         1
14 Two_Story_PUD_1946_and_Newer           96
15 PUD_Multilevel_Split_Level_Foyer       14
16 Two_Family_conversion_All_Styles_and_Ages 43
```

Label encoded

```
recipe(Sale_Price ~ ., data = ames_train) %>%
  step_integer(MS_SubClass) %>%
  prep(ames_train) %>%
  bake(ames_train) %>%
  count(MS_SubClass)
```

```
# A tibble: 16 x 2
  MS_SubClass      n
```


	<dbl>	<int>
1	1	753
2	2	91
3	3	5
4	4	11
5	5	211
6	6	395
7	7	98
8	8	17
9	9	75
10	10	32
11	11	66
12	12	145
13	13	1
14	14	96
15	15	14
16	16	43

Examples of ordinal features:

```
ames_train %>% select(contains("Qual"))
```

A tibble: 2,053 x 6

	Overall_Qual	Exter_Qual	Bsmt_Qual	Low_Qual_Fin_SF	Kitchen_Qual	Garage_Qual
	<fct>	<fct>	<fct>	<int>	<fct>	<fct>
1	Above_Average	Typical	Typical	0	Typical	Typical
2	Average	Typical	Typical	0	Typical	Typical
3	Above_Average	Typical	Typical	0	Good	Typical
4	Above_Average	Typical	Typical	0	Good	Typical
5	Very_Good	Good	Good	0	Good	Typical
6	Very_Good	Good	Good	0	Good	Typical
7	Good	Typical	Typical	0	Good	Typical
8	Above_Average	Typical	Good	0	Typical	Typical
9	Above_Average	Typical	Good	0	Typical	Typical
10	Good	Typical	Good	0	Good	Typical

... with 2,043 more rows

```
count(ames_train, Overall_Qual)
```

A tibble: 10 x 2

	Overall_Qual	n
	<fct>	<int>
1	Very_Poor	4
2	Poor	8
3	Fair	26
4	Below_Average	170
5	Average	564

```

6 Above_Average      511
7 Good               439
8 Very_Good          231
9 Excellent           77
10 Very_Excellent     23

```

```

# Label encoded
recipe(Sale_Price ~., data = ames_train) %>%
  step_integer(Overall_Qual) %>%
  prep(ames_train) %>%
  bake(ames_train) %>%
  count(Overall_Qual)

```

```

# A tibble: 10 x 2
  Overall_Qual      n
      <dbl> <int>
1           1      4
2           2      8
3           3     26
4           4    170
5           5   564
6           6   511
7           7   439
8           8   231
9           9    77
10          10    23

```

Alternatives

Target encoding:

```

ames_train %>%
  group_by(Neighborhood) %>%
  summarize(`Avg Sale_Price` = mean(Sale_Price, na.rm = TRUE)) %>%
  head(10) %>%
  kable(caption = "Example of target encoding the Neighborhood feature of the Ames housing data",
        kable_styling(bootstrap_options = "striped", full_width = TRUE))

```

```

ames_train %>%
  count(Neighborhood) %>%
  mutate(Proportion = n / sum(n)) %>%
  select(-n) %>%
  head(10) %>%
  kable(caption = "Example of categorical proportion encoding the Neighborhood feature of the Ames housing data",
        kable_styling(bootstrap_options = "striped", full_width = TRUE))

```

Table 1: Example of target encoding the Neighborhood feature of the Ames housing data set.

Neighborhood	Avg Sale_Price
North_Ames	144562.7
College_Creek	199831.7
Old_Town	122736.7
Edwards	130652.2
Somerset	227379.6
Northridge_Heights	323289.5
Gilbert	192162.9
Sawyer	136460.6
Northwest_Ames	187328.2
Sawyer_West	188644.6

Table 2: Example of categorical proportion encoding the Neighborhood feature of the Ames housing data set.

Neighborhood	Proportion
North_Ames	0.1451534
College_Creek	0.0910862
Old_Town	0.0832927
Edwards	0.0711154
Somerset	0.0623478
Northridge_Heights	0.0560156
Gilbert	0.0565027
Sawyer	0.0496834
Northwest_Ames	0.0467608
Sawyer_West	0.0414028

Dimension Reduction

Example PCA using *resample* package.

Data Recipe

Inputs:

	role	#variables
outcome		1
predictor		80

Operations:

Centering for all_numeric
Scaling for all_numeric
No PCA components were extracted.

Full Recipe

Full blueprint recipe applied to training and test data.

```
blueprint <- recipe(Sale_Price ~ ., data = ames_train) %>%  
  step_nzv(all_nominal()) %>%  
  step_integer(matches("Qual|Cond|QC|Qu")) %>%  
  step_center(all_numeric(), -all_outcomes()) %>%  
  step_scale(all_numeric(), -all_outcomes()) %>%  
  step_pca(all_numeric(), -all_outcomes())
```

blueprint

Data Recipe

Inputs:

	role	#variables
outcome		1
predictor		80

Operations:

Sparse, unbalanced variable filter on all_nominal
Integer encoding for matches, Qual|Cond|QC|Qu
Centering for all_numeric, -, all_outcomes()

Scaling for all_numeric, -, all_outcomes()
 No PCA components were extracted.

```
prepare <- prep(blueprint, training = ames_train)
prepare
```

Data Recipe

Inputs:

```
      role #variables
outcome      1
predictor    80
```

Training data contained 2053 data points and no missing data.

Operations:

Sparse, unbalanced variable filter removed Street, Alley, Land_Contour, ... [trained]
 Integer encoding for Condition_1, Overall_Qual, Overall_Cond, ... [trained]
 Centering for Lot_Frontage, Lot_Area, ... [trained]
 Scaling for Lot_Frontage, Lot_Area, ... [trained]
 PCA extraction with Lot_Frontage, Lot_Area, ... [trained]

```
baked_train <- bake(prepare, new_data = ames_train)
baked_test  <- bake(prepare, new_data = ames_test)
```

baked_train

A tibble: 2,053 x 27

	MS_SubClass	MS_Zoning	Lot_Shape	Lot_Config	Neighborhood	Bldg_Type	House_Style
	<fct>	<fct>	<fct>	<fct>	<fct>	<fct>	<fct>
1	One_Story_...	Resident...	Slightly...	Corner	North_Ames	OneFam	One_Story
2	One_Story_...	Resident...	Regular	Inside	North_Ames	OneFam	One_Story
3	One_Story_...	Resident...	Slightly...	Corner	North_Ames	OneFam	One_Story
4	Two_Story_...	Resident...	Slightly...	Inside	Gilbert	OneFam	Two_Story
5	One_Story_...	Resident...	Regular	Inside	Stone_Brook	TwnhsE	One_Story
6	One_Story_...	Resident...	Slightly...	Inside	Stone_Brook	TwnhsE	One_Story
7	Two_Story_...	Resident...	Regular	Inside	Gilbert	OneFam	Two_Story
8	Two_Story_...	Resident...	Slightly...	Corner	Gilbert	OneFam	Two_Story
9	Two_Story_...	Resident...	Slightly...	Inside	Gilbert	OneFam	Two_Story
10	One_Story_...	Resident...	Regular	Inside	Gilbert	OneFam	One_Story

... with 2,043 more rows, and 20 more variables: Roof_Style <fct>,
 # Exterior_1st <fct>, Exterior_2nd <fct>, Mas_Vnr_Type <fct>,
 # Foundation <fct>, Bsmt_Exposure <fct>, BsmtFin_Type_1 <fct>,
 # Central_Air <fct>, Electrical <fct>, Garage_Type <fct>,

```
# Garage_Finish <fct>, Paved_Drive <fct>, Fence <fct>, Sale_Type <fct>,
# Sale_Price <int>, PC1 <dbl>, PC2 <dbl>, PC3 <dbl>, PC4 <dbl>, PC5 <dbl>
```

Full recipe with cross-validation & grid search using carat.

k-Nearest Neighbors

```
2053 samples
  80 predictor
```

Recipe steps: nzv, integer, center, scale, dummy

Resampling: Cross-Validated (10 fold, repeated 5 times)

Summary of sample sizes: 1848, 1846, 1848, 1847, 1848, 1849, ...

Resampling results across tuning parameters:

k	RMSE	Rsquared	MAE
2	36116.20	0.8024564	22617.83
3	35077.80	0.8149018	21698.15
4	34670.83	0.8205474	21274.44
5	34240.24	0.8280936	21067.36
6	33813.06	0.8346543	20889.18
7	33517.75	0.8404839	20777.87
8	33324.82	0.8440427	20647.95
9	33148.42	0.8468769	20598.96
10	33059.04	0.8488142	20610.92
11	33048.49	0.8500846	20657.02
12	32951.81	0.8521235	20680.45
13	33017.22	0.8529401	20736.74
14	32959.37	0.8548174	20753.98
15	32976.91	0.8558820	20797.87
16	33013.07	0.8565768	20849.17
17	33028.05	0.8571048	20907.07
18	33119.22	0.8568609	21015.28
19	33124.87	0.8572290	21065.00
20	33199.35	0.8571015	21147.53
21	33274.67	0.8566490	21204.72
22	33291.52	0.8568531	21250.21
23	33323.99	0.8571500	21295.17
24	33399.65	0.8570289	21366.10
25	33470.78	0.8569297	21439.81

RMSE was used to select the optimal model using the smallest value.

The final value used for the model was k = 12.

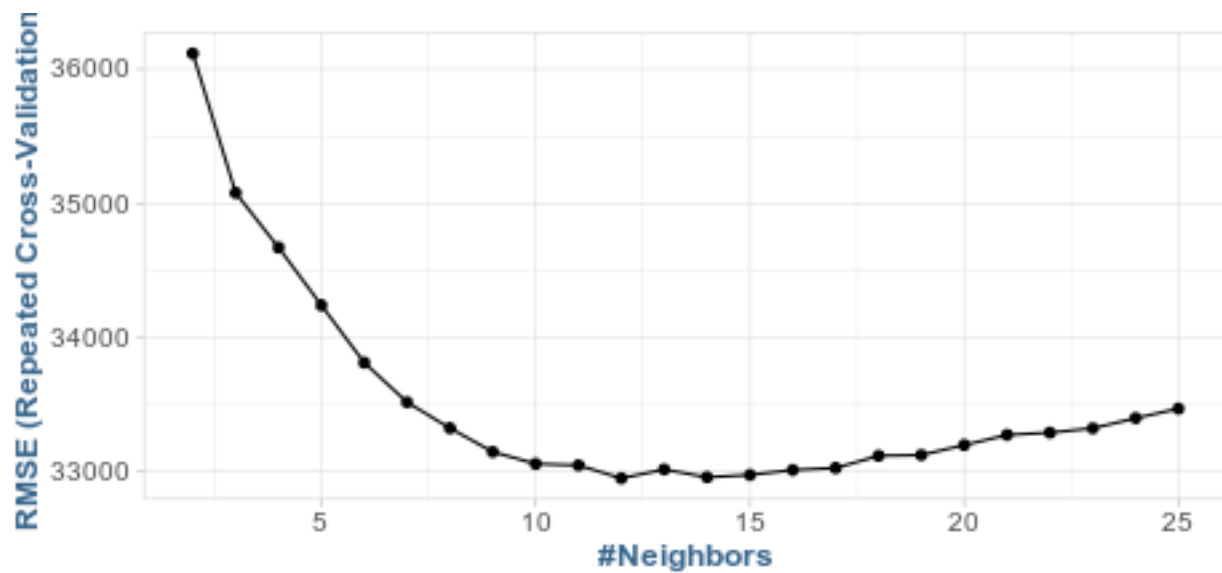


Figure 5: Results from the same grid search performed in Section 2.7 but with feature engineering performed within each resample.