

Feature and Target Engineering

Data Set

h2o

```
ames <- AmesHousing::make_ames()
ames.h2o <- as.h2o(ames)
```

stratified (*Sale_Price*) training sample

```
set.seed(123)

split <- initial_split(ames, prop = 0.7,
                       strata = "Sale_Price")

ames_train <- training(split)
ames_test <- testing(split)
```

log transformation (*Sale_Price*)

```
ames_recipe <- recipe(Sale_Price ~ ., data = ames_train) %>%
  step_log(all_outcomes())

ames_recipe
```

Data Recipe

Inputs:

	role	#variables
outcome		1
predictor		80

Operations:

Log transformation on *all_outcomes*

Box-Cox transformation (example)

```
lambda <- 3

y <- forecast::BoxCox(10, lambda)

inv_box_cox <- function(x, lambda) {
  # for Box-Cox, lambda = 0 -> log transform
  if(lambda == 0) exp(x) else (lambda*x + 1)^(1/lambda)
```

```

}

inv_box_cox(y, lambda)

[1] 10
attr(,"lambda")
[1] 3

# Log transformation
train_log_y <- log(ames_train$Sale_Price)
test_log_y  <- log(ames_train$Sale_Price)

# Box Cox transformation
lambda <- forecast::BoxCox.lambda(ames_train$Sale_Price)
train_bc_y <- forecast::BoxCox(ames_train$Sale_Price, lambda)
test_bc_y  <- forecast::BoxCox(ames_test$Sale_Price, lambda)

# Plot differences
levs <- c("Normal", "Log_Transform", "BoxCox_Transform")
data.frame(
  Normal = ames_train$Sale_Price,
  Log_Transform = train_log_y,
  BoxCox_Transform = train_bc_y
) %>%
  gather(Transform, Value) %>%
  mutate(Transform = factor(Transform, levels = levs)) %>%
  ggplot(aes(Value, fill = Transform)) +
    geom_histogram(show.legend = FALSE, bins = 40) +
    facet_wrap(~ Transform, scales = "free_x")

```





Figure 1: Response variable transformations.

Missing Values

```
sum(is.na(AmesHousing::ames_raw))
```

```
[1] 13997
```

```
AmesHousing::ames_raw %>%
  is.na() %>%
  reshape2::melt() %>%
  ggplot(aes(Var2, Var1, fill = value)) +
    geom_raster() +
    coord_flip() +
    scale_y_continuous(NULL, expand = c(0,0)) +
    scale_fill_grey(name = "",
                    labels = c("Present",
                              "Missing")) +
  xlab("Observations") +
  theme(axis.text.y = element_text(size = 4))
```



Missing Garage?

```
AmesHousing::ames_raw %>%
  filter(is.na(`Garage Type`)) %>%
  select(starts_with("Garage"))
```

```
# A tibble: 157 x 7
```

```
  `Garage Type` `Garage Yr Blt` `Garage Finish` `Garage Cars` `Garage Area`
  <chr>          <int> <chr>          <int>          <int>
```


```

1 <NA>      NA <NA>      0      0
2 <NA>      NA <NA>      0      0
3 <NA>      NA <NA>      0      0
4 <NA>      NA <NA>      0      0
5 <NA>      NA <NA>      0      0
6 <NA>      NA <NA>      0      0
7 <NA>      NA <NA>      0      0
8 <NA>      NA <NA>      0      0
9 <NA>      NA <NA>      0      0
10 <NA>     NA <NA>      0      0
# ... with 147 more rows, and 2 more variables: `Garage Qual` <chr>, `Garage
#   Cond` <chr>

```

Missing values w/cluster (*visdat*)

```
vis_miss(AmesHousing::ames_raw, cluster = T)
```



Missing Value Imputation

basic descriptive statistic

```
ames_recipe %>%
  step_medianimpute(Gr_Liv_Area)
```

Data Recipe

Inputs:

```

    role #variables
outcome          1
predictor        80

```

Operations:

Log transformation on all_outcomes

Median Imputation for Gr_Liv_Area

KNN approach (typical k = 5-10)

```

ames_recipe %>%
  step_knnimpute(all_predictors(), neighbors = 6)

```

Data Recipe

Inputs:

```

    role #variables
outcome          1
predictor        80

```

Operations:

Log transformation on all_outcomes

K-nearest neighbor imputation for all_predictors

```

impute_ames <- ames_train
set.seed(123)
index <- sample(seq_along(impute_ames$Gr_Liv_Area), 50)
actuals <- ames_train[index, ]
impute_ames$Gr_Liv_Area[index] <- NA

p1 <- ggplot() +
  geom_point(data = impute_ames, aes(Gr_Liv_Area, Sale_Price), alpha = .2) +
  geom_point(data = actuals, aes(Gr_Liv_Area, Sale_Price), color = "red") +
  scale_x_log10(limits = c(300, 5000)) +
  scale_y_log10(limits = c(10000, 500000)) +
  ggtitle("Actual values")

# Mean imputation
mean_juiced <- recipe(Sale_Price ~ ., data = impute_ames) %>%
  step_meanimpute(Gr_Liv_Area) %>%
  prep(training = impute_ames, retain = TRUE) %>%

```

```

juice()
mean_impute <- mean_juiced[index, ]

p2 <- ggplot() +
  geom_point(data = actuals, aes(Gr_Liv_Area, Sale_Price), color = "red") +
  geom_point(data = mean_impute, aes(Gr_Liv_Area, Sale_Price), color = "blue") +
  scale_x_log10(limits = c(300, 5000)) +
  scale_y_log10(limits = c(10000, 500000)) +
  ggtitle("Mean Imputation")

# KNN imputation
knn_juiced <- recipe(Sale_Price ~ ., data = impute_ames) %>%
  step_knnimpute(Gr_Liv_Area) %>%
  prep(training = impute_ames, retain = TRUE) %>%
  juice()
knn_impute <- knn_juiced[index, ]

p3 <- ggplot() +
  geom_point(data = actuals, aes(Gr_Liv_Area, Sale_Price), color = "red") +
  geom_point(data = knn_impute, aes(Gr_Liv_Area, Sale_Price), color = "blue") +
  scale_x_log10(limits = c(300, 5000)) +
  scale_y_log10(limits = c(10000, 500000)) +
  ggtitle("KNN Imputation")

# Bagged imputation
bagged_juiced <- recipe(Sale_Price ~ ., data = impute_ames) %>%
  step_bagimpute(Gr_Liv_Area) %>%
  prep(training = impute_ames, retain = TRUE) %>%
  juice()
bagged_impute <- bagged_juiced[index, ]

p4 <- ggplot() +
  geom_point(data = actuals, aes(Gr_Liv_Area, Sale_Price), color = "red") +
  geom_point(data = bagged_impute, aes(Gr_Liv_Area, Sale_Price), color = "blue") +
  scale_x_log10(limits = c(300, 5000)) +
  scale_y_log10(limits = c(10000, 500000)) +
  ggtitle("Bagged Trees Imputation")

gridExtra::grid.arrange(p1, p2, p3, p4, nrow = 2)

```

Warning: Removed 63 rows containing missing values (geom_point).

Increase in training time by model type:

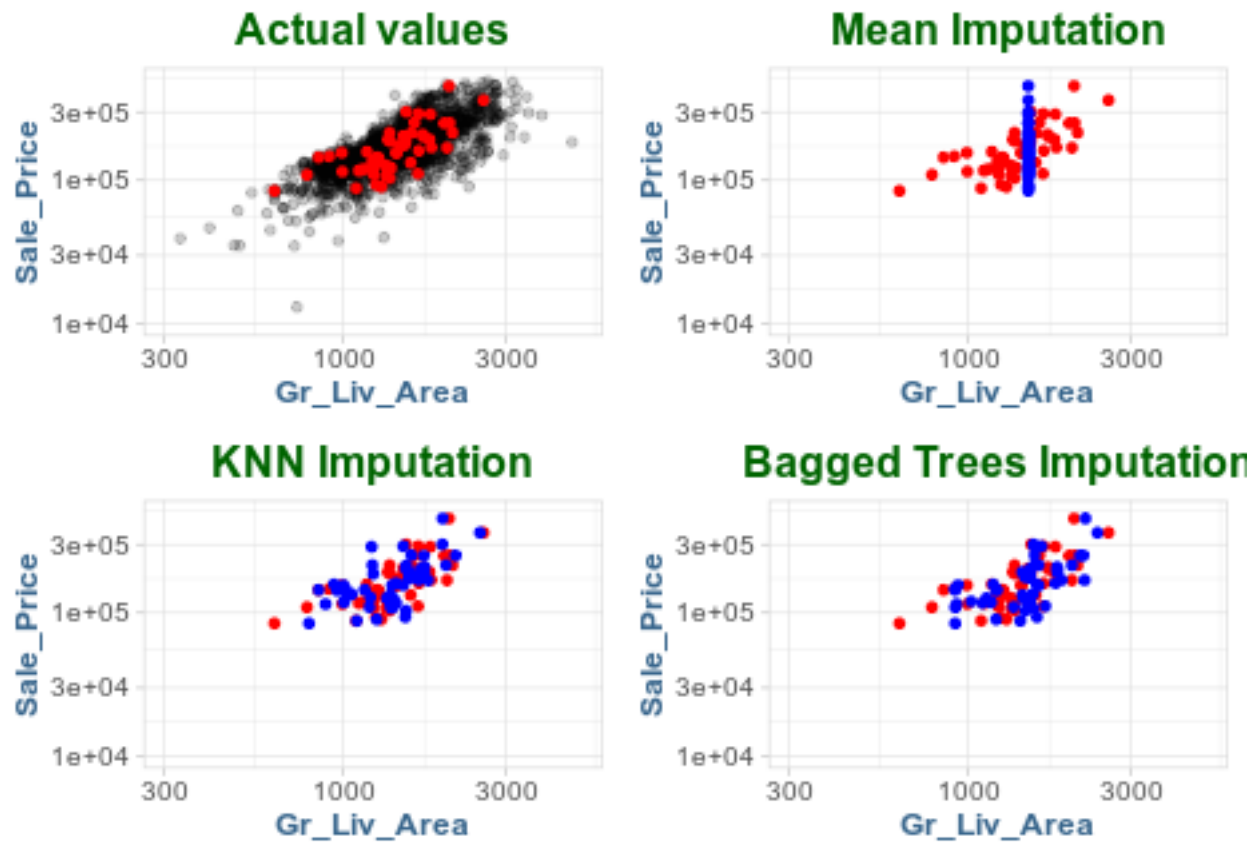


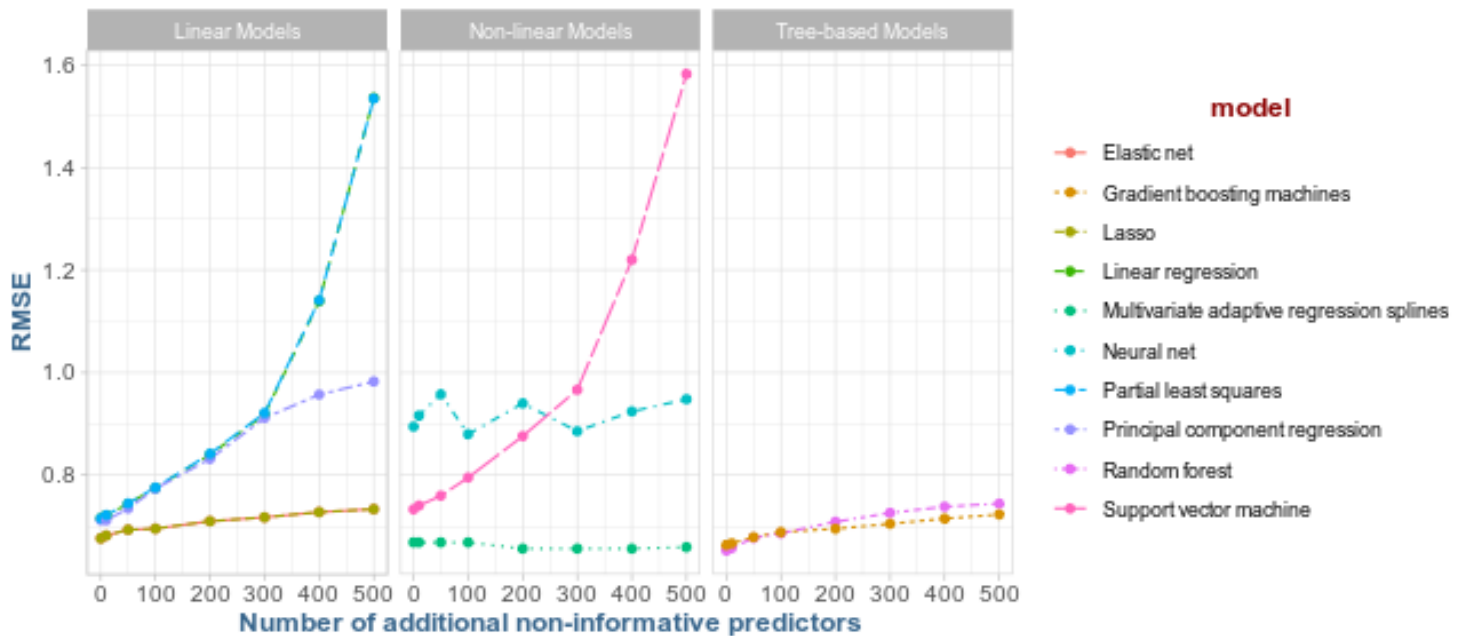
Figure 2: Comparison of three different imputation methods. The red points represent actual values which were removed and made missing and the blue points represent the imputed values. Estimated statistic imputation methods (i.e. mean, median) merely predict the same value for each observation and can reduce the signal between a feature and the response; whereas KNN and tree-based procedures tend to maintain the feature distribution and relationship.

```
model_results <- read_csv(paste0(data.dir, "feature-selection-impacts-results.csv")) %>%
  mutate(type = case_when(
    model %in% c("lm", "pcr", "pls", "glmnet", "lasso") ~ "Linear Models",
    model %in% c("earth", "svmLinear", "nn") ~ "Non-linear Models",
    TRUE ~ "Tree-based Models"
  )) %>%
  mutate(model = case_when(
    model == "lm" ~ "Linear regression",
    model == "earth" ~ "Multivariate adaptive regression splines",
    model == "gbm" ~ "Gradient boosting machines",
    model == "glmnet" ~ "Elastic net",
    model == "lasso" ~ "Lasso",
    model == "nn" ~ "Neural net",
    model == "pcr" ~ "Principal component regression",
    model == "pls" ~ "Partial least squares",
    model == "ranger" ~ "Random forest",
    TRUE ~ "Support vector machine"
  ))
```

Parsed with column specification:

```
cols(
  model = col_character(),
  NIP = col_double(),
  RMSE = col_double(),
  time = col_double()
)

ggplot(model_results, aes(NIP, RMSE, color = model, lty = model)) +
  geom_line() +
  geom_point() +
  facet_wrap(~ type, nrow = 1) +
  xlab("Number of additional non-informative predictors")
```

```
model_results %>%
  group_by(model) %>%
  mutate(
    time_impact = time / first(time),
    time_impact = time_impact - 1
  ) %>%
  ggplot(aes(NIP, time_impact, color = model, lty = model)) +
  geom_line() +
  geom_point() +
  facet_wrap(~ type, nrow = 1) +
  scale_y_continuous("Percent increase in training time",
    labels = scales::percent) +
  xlab("Number of additional non-informative predictors")
```

Rules of thumb for zero variance features:

- The fraction of unique values over the sample size is low (say < 10%)
- The ratio of the frequency of the most prevalent value to the frequency of the second most prevalent value is large (say > 20%)

If both of these criteria are met, then it is often advantageous to remove them from the model.

```
caret::nearZeroVar(ames_train, saveMetrics = T) %>%
  rownames_to_column() %>%
  filter(nzv)
```

	rowname	freqRatio	percentUnique	zeroVar	nzv
1	Street	292.28571	0.09741841	FALSE	TRUE
2	Alley	20.52688	0.14612762	FALSE	TRUE

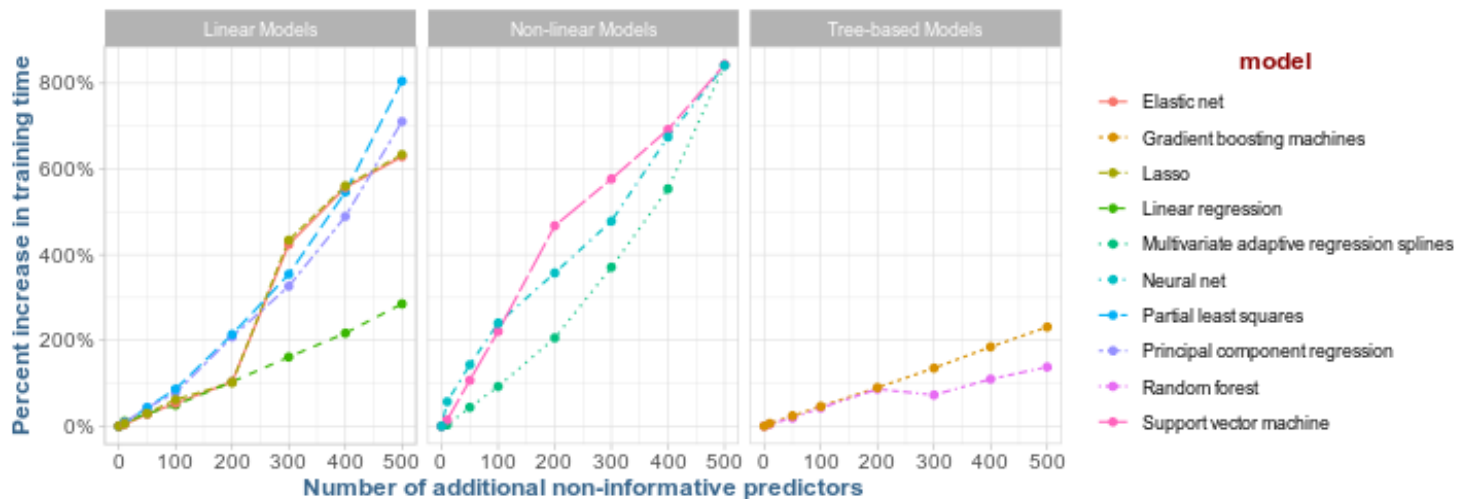


Figure 3: Impact in model training time as non-informative predictors are added.

3	Land_Contour	22.28916	0.19483682	FALSE	TRUE
4	Utilities	1025.00000	0.14612762	FALSE	TRUE
5	Land_Slope	22.76744	0.14612762	FALSE	TRUE
6	Condition_2	203.10000	0.34096444	FALSE	TRUE
7	Roof_Mat1	126.50000	0.24354603	FALSE	TRUE
8	Bsmt_Cond	19.93478	0.29225524	FALSE	TRUE
9	BsmtFin_Type_2	21.50617	0.34096444	FALSE	TRUE
10	Heating	101.05000	0.24354603	FALSE	TRUE
11	Low_Qual_Fin_SF	1013.00000	1.31514856	FALSE	TRUE
12	Kitchen_AbvGr	23.68675	0.19483682	FALSE	TRUE
13	Functional	38.18000	0.34096444	FALSE	TRUE
14	Enclosed_Porch	100.94118	7.40379932	FALSE	TRUE
15	Three_season_porch	674.66667	1.16902094	FALSE	TRUE
16	Screen_Porch	234.87500	4.52995616	FALSE	TRUE
17	Pool_Area	2045.00000	0.43838285	FALSE	TRUE
18	Pool_QC	681.66667	0.24354603	FALSE	TRUE
19	Misc_Feature	30.49231	0.19483682	FALSE	TRUE
20	Misc_Val	165.33333	1.41256698	FALSE	TRUE

Numeric Feature Engineering

Skewness can have a drastic impact on the performance of GLMs & regularized models.

Non-parametric models are rarely affected by skewed features; however, normalizing features will not have a negative effect on these models' performance. For example, normalizing features will only shift the optimal split points in tree-based algorithms. Consequently, when in doubt, normalize.

Skewness

Data Recipe

Inputs:

	role	#variables
outcome		1
predictor		80

Operations:

Yeo-Johnson transformation on all_numeric

Standardization

```
ames_recipe %>%  
  step_center(all_numeric(), -all_outcomes()) %>%  
  step_scale(all_numeric(), -all_outcomes())
```

Data Recipe

Inputs:

	role	#variables
outcome		1
predictor		80

Operations:

Log transformation on all_outcomes

Centering for all_numeric, -, all_outcomes()

Scaling for all_numeric, -, all_outcomes()

```
set.seed(123)  
x1 <- tibble(  
  variable = "x1",  
  `Real Value` = runif(25, min = -30, max = 5),  
  `Standardized Value` = scale(`Real Value`) %>% as.numeric()  
)  
  
set.seed(456)  
x2 <- tibble(  
  variable = "x1",  
  `Real Value` = runif(25, min = -30, max = 5),  
  `Standardized Value` = scale(`Real Value`) %>% as.numeric()  
)
```

```

variable = "x2",
`Real value` = rlnorm(25, log(25)),
`Standardized value` = scale(`Real value`) %>% as.numeric()
)

set.seed(789)
x3 <- tibble(
  variable = "x3",
  `Real value` = rnorm(25, 150, 15),
  `Standardized value` = scale(`Real value`) %>% as.numeric()
)

x1 %>%
  bind_rows(x2) %>%
  bind_rows(x3) %>%
  gather(key, value, -variable) %>%
  mutate(variable = factor(variable, levels = c("x3", "x2", "x1"))) %>%
  ggplot(aes(value, variable)) +
    geom_point(alpha = .6) +
    facet_wrap(~ key, scales = "free_x") +
    ylab("Feature") +
    xlab("Value")

```

Warning: Removed 150 rows containing missing values (geom_point).




Figure 4: Standardizing features allows all features to be compared on a common value scale regardless of their real value differences.

Categorical Feature Engineering

Lumping

When a feature contains levels that have few observations.

For example:

```
count(ames_train, Neighborhood) %>% arrange(n)
```

```
# A tibble: 27 x 2
  Neighborhood      n
  <fct>          <int>
1 Green_Hills      2
2 Greens           7
3 Blueste          8
4 Northpark_Villa 17
5 Briardale        18
6 Veenker          20
7 Bloomington_Heights 21
8 South_and_West_of_Iowa_State_University 27
9 Meadow_Village   29
10 Clear_Creek     31
# ... with 17 more rows
```

```
count(ames_train, Screen_Porch) %>% arrange(n)
```

```
# A tibble: 93 x 2
  Screen_Porch      n
  <int> <int>
1      40      1
2      63      1
3      80      1
4      92      1
5      94      1
6      99      1
7     104      1
8     109      1
9     110      1
10     111      1
# ... with 83 more rows
```

We can benefit from lumping these together into an “other” category when they contain less than 10% of the training sample.

Note: This can have an adverse effect on performance

```
lumping <- recipe(Sale_Price ~., data = ames_train) %>%
  step_other(Neighborhood, threshold = 0.01,
             other = "other") %>%
  step_other(Screen_Porch, threshold = 0.1,
             other = ">0")

apply_2_training <- prep(lumping, training = ames_train) %>%
  bake(ames_train)

# New distribution of Neighborhood
count(apply_2_training, Neighborhood) %>% arrange(n)
```

```
# A tibble: 22 x 2
  Neighborhood      n
  <fct>           <int>
1 Bloomington_Heights 21
2 South_and_West_of_Iowa_State_University 27
3 Meadow_Village      29
4 Clear_Creek         31
5 Stone_Brook         34
6 Northridge         48
7 Timberland         55
8 Iowa_DOT_and_Rail_Road 62
9 Crawford           72
10 other             72
# ... with 12 more rows
```

```
# New distribution of Screen_Porch
count(apply_2_training, Screen_Porch) %>% arrange(n)
```

```
# A tibble: 2 x 2
  Screen_Porch      n
  <fct>         <int>
1 >0           174
2 0           1879
```

```
dat <- data.table(id = 1:9, x = rep(c("a", "b", "c"), 3))
dat
```

```
id x
1: 1 a
2: 2 b
3: 3 c
4: 4 a
5: 5 b
6: 6 c
```

```
7: 7 a
8: 8 b
9: 9 c
```

```
# full-rank
```

```
dat[, .(id,
        `X = a` = as.numeric(x == "a"),
        `X = b` = as.numeric(x == "b"),
        `X = c` = as.numeric(x == "c")) ]
```

	id	X = a	X = b	X = c
1:	1	1	0	0
2:	2	0	1	0
3:	3	0	0	1
4:	4	1	0	0
5:	5	0	1	0
6:	6	0	0	1
7:	7	1	0	0
8:	8	0	1	0
9:	9	0	0	1

```
# one-hot (leave one out)
```

```
dat[, .(id,
        `X = a` = as.numeric(x == "a"),
        `X = b` = as.numeric(x == "b")) ]
```

	id	X = a	X = b
1:	1	1	0
2:	2	0	1
3:	3	0	0
4:	4	1	0
5:	5	0	1
6:	6	0	0
7:	7	1	0
8:	8	0	1
9:	9	0	0