

## Linear Regression

### Data Set

h2o

```
ames <- AmesHousing::make_ames()
ames.h2o <- as.h2o(ames)
```

stratified (*Sale\_Price*) training sample

```
set.seed(123)

split <- initial_split(ames, prop = 0.7,
                       strata = "Sale_Price")

ames_train <- training(split)
ames_test <- testing(split)
```

### Simple Linear Model

```
model1 <- lm(Sale_Price ~ Gr_Liv_Area, data = ames_train)
```

*# Fitted regression line (full training)*


```
p1 <- model1 %>%
  broom::augment() %>%
  ggplot(aes(Gr_Liv_Area, Sale_Price)) +
  geom_point(size = 1, alpha = 0.3) +
  geom_smooth(se = F, method = "lm") +
  scale_y_continuous(labels = scales::dollar) +
  ggtitle("Fitted regression line")
```

*# Fitted regression line (restricted range)*


```
p2 <- model1 %>%
  broom::augment() %>%
  ggplot(aes(Gr_Liv_Area, Sale_Price)) +
  geom_segment(aes(x = Gr_Liv_Area, y = Sale_Price,
                  xend = Gr_Liv_Area, yend = .fitted),
              alpha = .3) +
  geom_point(size = 1, alpha = 0.3) +
  geom_smooth(se = F, method = "lm") +
  scale_y_continuous(labels = scales::dollar) +
  ggtitle("Fitted regression line (with residuals)")
```

```
grid.arrange(p1, p2, nrow = 2)
```

Fitted regression line



Fitted regression line (with residuals)



```
summary(model1)
```

Call:

```
lm(formula = Sale_Price ~ Gr_Liv_Area, data = ames_train)
```

Residuals:

Min	1Q	Median	3Q	Max
-361143	-30668	-2449	22838	331357

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	8732.938	3996.613	2.185	0.029 *
Gr_Liv_Area	114.876	2.531	45.385	<0.0000000000000002 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 56700 on 2051 degrees of freedom

Multiple R-squared: 0.5011, Adjusted R-squared: 0.5008

F-statistic: 2060 on 1 and 2051 DF, p-value: < 0.00000000000000022

```
[1] 56704.78
```

[1] 3215432370

## Inference

The variability of an estimate is its *standard error (SE)*, the square root of its variance.

t-test for the coefficients are simply the estimated coefficient divided by the standard error (t value = Estimate / Std. Error)

t-test measure the number of standard deviations each coefficient is away from zero (basically  $\text{abs}(T) > 2$  is significant at 95% conf)

The confidence interval for coefficients is:

$$\hat{\beta}_j \pm t_{1-\alpha/2, n-p} \hat{SE}(\hat{\beta}_j)$$

```
confint(model1, level = .95)
```

```

                2.5 %      97.5 %
(Intercept) 895.0961 16570.7805
Gr_Liv_Area 109.9121  119.8399
```

Interpretation: We are 95% confident that each one unit increase in Gr\_Liv\_Area adds between 109.9 and 119.8 dollars to the sale price.

Linear Regression Assumptions:

- 1.) Independent observations
- 2.) The random errors have mean zero, and constant variance
- 3.) The random errors are normally distributed

## Multiple Linear Regression

```
(model2 <- lm(Sale_Price ~ Gr_Liv_Area + Year_Built, data = ames_train))
```

Call:

```
lm(formula = Sale_Price ~ Gr_Liv_Area + Year_Built, data = ames_train)
```

Coefficients:

```

(Intercept)  Gr_Liv_Area  Year_Built
-2123054.21      99.18      1093.48
```

*# Equivalent*

```
(model2 <- update(model1, . ~ . + Year_Built))
```

Call:

```
lm(formula = Sale_Price ~ Gr_Liv_Area + Year_Built, data = ames_train)
```

Coefficients:

```
(Intercept)  Gr_Liv_Area  Year_Built
-2123054.21      99.18      1093.48
```

```
round(coef(model2), 3)
```

```
(Intercept)  Gr_Liv_Area  Year_Built
-2123054.207      99.176      1093.485
```

```
summary(model3 <- lm(Sale_Price ~ Gr_Liv_Area + Year_Built + Gr_Liv_Area:Year_Built, data = ames_train))
```

Call:

```
lm(formula = Sale_Price ~ Gr_Liv_Area + Year_Built + Gr_Liv_Area:Year_Built,
    data = ames_train)
```

Residuals:

```
      Min       1Q   Median       3Q      Max
-440543  -25191   -1896    17599   281542
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	382194.30149	209192.64043	1.827	0.0678
Gr_Liv_Area	-1483.88103	125.65052	-11.810	<0.0000000000000002
Year_Built	-179.79795	106.40443	-1.690	0.0912
Gr_Liv_Area:Year_Built	0.80371	0.06378	12.601	<0.0000000000000002

```
(Intercept)      .
Gr_Liv_Area      ***
Year_Built       .
Gr_Liv_Area:Year_Built ***
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Residual standard error: 44810 on 2049 degrees of freedom

Multiple R-squared: 0.6887, Adjusted R-squared: 0.6883

F-statistic: 1511 on 3 and 2049 DF, p-value: < 0.00000000000000022

```
round(coef(model3), 3)
```

```
(Intercept)      Gr_Liv_Area      Year_Built
382194.301      -1483.881      -179.798
Gr_Liv_Area:Year_Built
```

0.804

```
fit1 <- lm(Sale_Price ~ Gr_Liv_Area + Year_Built, data = ames_train)
fit2 <- lm(Sale_Price ~ Gr_Liv_Area * Year_Built, data = ames_train)

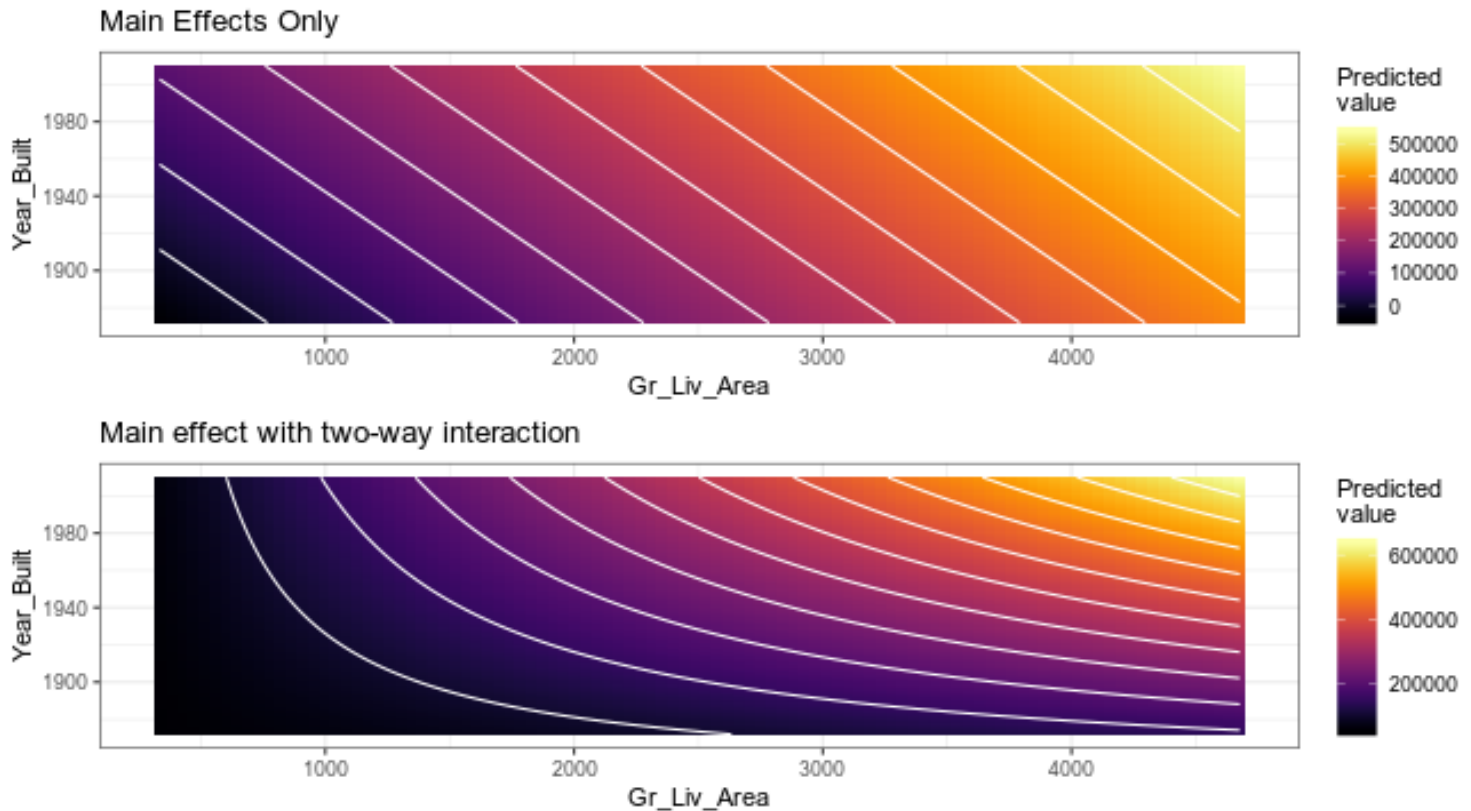
# Regression plane
plot_grid <- expand_grid(
  Gr_Liv_Area = seq(from = min(ames_train$Gr_Liv_Area), to = max(ames_train$Gr_Liv_Area),
    length = 100),
  Year_Built = seq(from = min(ames_train$Year_Built), to = max(ames_train$Year_Built),
    length = 100)
)

plot_grid$y1 <- predict(fit1, newdata = plot_grid)
plot_grid$y2 <- predict(fit2, newdata = plot_grid)

# Level plots
p1 <- ggplot(plot_grid, aes(x = Gr_Liv_Area, y = Year_Built,
  z = y1, fill = y1)) +
  geom_tile() +
  geom_contour(color = "white") +
  viridis::scale_fill_viridis(name = "Predicted\nvalue", option = "inferno") +
  theme_bw() +
  ggtitle("Main Effects Only")

p2 <- ggplot(plot_grid, aes(x = Gr_Liv_Area, y = Year_Built,
  z = y2, fill = y2)) +
  geom_tile() +
  geom_contour(color = "white") +
  viridis::scale_fill_viridis(name = "Predicted\nvalue", option = "inferno") +
  theme_bw() +
  ggtitle("Main effect with two-way interaction")

gridExtra::grid.arrange(p1, p2, nrow = 2)
```



### Full Model

```
model3 <- lm(Sale_Price ~ ., data = ames_train)
broom::tidy(model3)
```

# A tibble: 283 x 5

term	estimate	std.error	statistic	p.value
<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1 (Intercept)	-5.61e6	11261881.	-0.498	0.618
2 MS_SubClassOne_Story_1945_and_Older	3.56e3	3843.	0.926	0.355
3 MS_SubClassOne_Story_with_Finished_Atti...	1.28e4	12834.	0.997	0.319
4 MS_SubClassOne_and_Half_Story_Unfinishe...	8.73e3	12871.	0.678	0.498
5 MS_SubClassOne_and_Half_Story_Finished_...	4.11e3	6226.	0.660	0.509
6 MS_SubClassTwo_Story_1946_and_Newer	-1.09e3	5790.	-0.189	0.850
7 MS_SubClassTwo_Story_1945_and_Older	7.14e3	6349.	1.12	0.261
8 MS_SubClassTwo_and_Half_Story_All_Ages	-1.39e4	11003.	-1.27	0.206
9 MS_SubClassSplit_or_Multilevel	-1.15e4	10512.	-1.09	0.276
10 MS_SubClassSplit_Foyer	-4.39e3	8057.	-0.545	0.586

# ... with 273 more rows

### Assessing Model Accuracy

Models 1/2/3:

```
# Train model using 10-fold cross-validation

set.seed(123) # for reproducibility

(cv_model <- train(
  form = Sale_Price ~ Gr_Liv_Area,
  data = ames_train,
  method = "lm",
  trControl = trainControl(method = "cv", number = 10)
))
```

Linear Regression

2053 samples  
1 predictor

No pre-processing

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 1846, 1848, 1848, 1848, 1848, 1848, ...

Resampling results:

RMSE	Rsquared	MAE
56410.89	0.5069425	39169.09

Tuning parameter 'intercept' was held constant at a value of TRUE

```
set.seed(123)

cv_model2 <- train(
  Sale_Price ~ Gr_Liv_Area + Year_Built,
  data = ames_train,
  method = "lm",
  trControl = trainControl(method = "cv", number = 10)
)

set.seed(123)

suppressWarnings({
  cv_model3 <- train(
    Sale_Price ~ .,
    data = ames_train,
    method = "lm",
    trControl = trainControl(method = "cv", number = 10)
  )
})
```

**Accuracy:**

Call:

```
summary.resamples(object = resamples(list(model1 = cv_model, model2
= cv_model2, model3 = cv_model3)))
```

Models: model1, model2, model3

Number of resamples: 10

**MAE**

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
model1	34457.58	36323.74	38943.81	39169.09	41660.81	45005.17	0
model2	28094.79	30594.47	31959.30	32246.86	34210.70	37441.82	0
model3	12458.27	15420.10	16484.77	16258.84	17262.39	19029.29	0

**RMSE**

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
model1	47211.34	52363.41	54948.96	56410.89	60672.31	67679.05	0
model2	37698.17	42607.11	45407.14	46292.38	49668.59	54692.06	0
model3	20844.33	22581.04	24947.45	26098.00	27695.65	39521.49	0

**Rsquared**

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
model1	0.3598237	0.4550791	0.5289068	0.5069425	0.5619841	0.5965793	0
model2	0.5714665	0.6392504	0.6800818	0.6703298	0.7067458	0.7348562	0
model3	0.7869022	0.9018567	0.9104351	0.8949642	0.9166564	0.9303504	0

**Model Concerns / Assumptions****1.) Linear Relationships**

Possible solution to non-linear relationship is by variable transformations:

```
p1 <- ggplot(ames_train, aes(Year_Built, Sale_Price)) +
  geom_point(size = 1, alpha = .4) +
  geom_smooth(se = F) +
  scale_y_continuous(labels = scales::dollar) +
  xlab("Year Built") +
  ggtitle(paste("Non-transformed variables with a \n", "non-linear relationship"))

p2 <- ggplot(ames_train, aes(Year_Built, Sale_Price)) +
  geom_point(size = 1, alpha = .4) +
  geom_smooth(se = F, method = "lm") +
  scale_y_log10("Sale Price", labels = scales::dollar, breaks = seq(0, 400000, by = 100000)) +
```




```
xlab("Year Built") +
ggtitle(paste("Transforming variables can provide a \n", "near-linear relationship"))


gridExtra::grid.arrange(p1, p2, nrow = 1)
```

`geom\_smooth()` using method = 'gam' and formula 'y ~ s(x, bs = "cs")'

**Non-transformed variables with a non-linear relationship**



**Transforming variables can provide near-linear relationship**



## 2.) Constant variance among residuals

Linear models assume constant variance among error terms (*homoscedasticity*).

Notice the cone shape in Model 1:



```
df1 <- broom::augment(cv_model$finalModel, data = ames_train)

p1 <- ggplot(df1, aes(.fitted, .resid)) +
  geom_point(size = 1, alpha = .4) +
  xlab("Predicted Values") +
  ylab("Residuals") +
  ggtitle("Model 1", subtitle = "Sale_Price ~ Gr_Liv_Area")

df2 <- broom::augment(cv_model3$finalModel, data = ames_train)

p2 <- ggplot(df2, aes(.fitted, .resid)) +
  geom_point(size = 1, alpha = .4) +
  xlab("Predicted Values") +
  ylab("Residuals") +
  ggtitle("Model 3", subtitle = "Sale_Price ~ .")
```

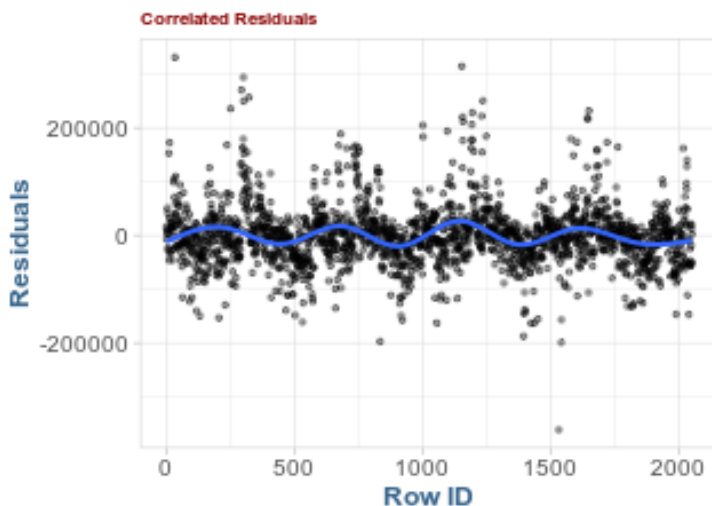
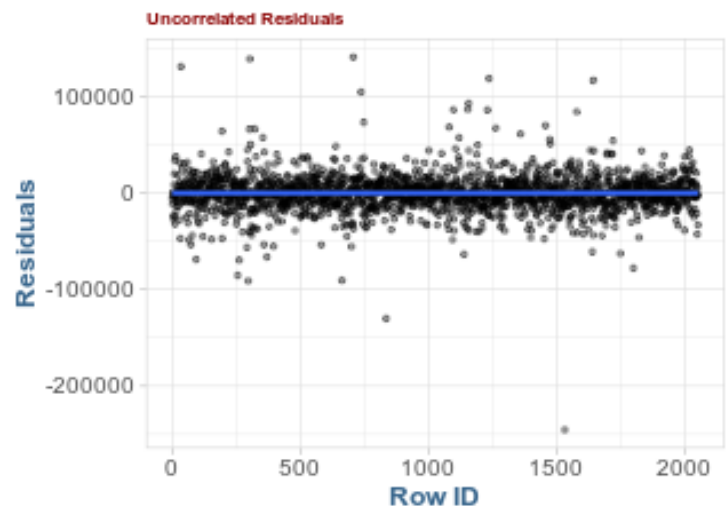
```
gridExtra::grid.arrange(p1, p2, nrow = 1)
```

**Model 1****Model 3**

### 3.) No autocorrelation

Residuals should be uncorrelated and i.i.d.

```
`geom_smooth()` using method = 'gam' and formula 'y ~ s(x, bs = "cs")'
`geom_smooth()` using method = 'gam' and formula 'y ~ s(x, bs = "cs")'
```

**Model 1****Model 3**

Model 1 has a distinct pattern to the residuals (due to being ordered by neighborhood, which is unaccounted for in the model)