

Chapter 4

Book Work

Simple Logistic Regression

```
data.default <- data.table(ISLR::Default)[, dflt := ifelse(default == "Yes", 1, 0)]

summary(model1 <- glm(dflt ~ balance, data = data.default, family = "binomial"))
```

Call:

```
glm(formula = dflt ~ balance, family = "binomial", data = data.default)
```

Deviance Residuals:

| Min | 1Q | Median | 3Q | Max |
|---------|---------|---------|---------|--------|
| -2.2697 | -0.1465 | -0.0589 | -0.0221 | 3.7589 |

Coefficients:

| | Estimate | Std. Error | z value | Pr(> z) |
|-------------|------------|------------|---------|------------|
| (Intercept) | -1.065e+01 | 3.612e-01 | -29.49 | <2e-16 *** |
| balance | 5.499e-03 | 2.204e-04 | 24.95 | <2e-16 *** |

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 2920.6 on 9999 degrees of freedom
 Residual deviance: 1596.5 on 9998 degrees of freedom
 AIC: 1600.5

Number of Fisher Scoring iterations: 8

```
predict(model1, newdata = data.frame( balance = c(1000, 2000) ), type = "response")
```

| 1 | 2 |
|-------------|-------------|
| 0.005752145 | 0.585769370 |

```
data.default[, is_student := ifelse(student == "Yes", 1, 0)]
```

```
summary(model2 <- glm(dflt ~ is_student, data = data.default, family = "binomial"))
```

Call:

```
glm(formula = dflt ~ is_student, family = "binomial", data = data.default)
```

Deviance Residuals:

| Min | 1Q | Median | 3Q | Max |
|---------|---------|---------|---------|--------|
| -0.2970 | -0.2970 | -0.2434 | -0.2434 | 2.6585 |

Coefficients:

| | Estimate | Std. Error | z value | Pr(> z) |
|-------------|----------|------------|---------|--------------|
| (Intercept) | -3.50413 | 0.07071 | -49.55 | < 2e-16 *** |
| is_student | 0.40489 | 0.11502 | 3.52 | 0.000431 *** |

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 2920.6 on 9999 degrees of freedom
 Residual deviance: 2908.7 on 9998 degrees of freedom
 AIC: 2912.7

Number of Fisher Scoring iterations: 6

```
predict(model2, newdata = data.frame( is_student = c(1, 0) ), type = "response")
```

| 1 | 2 |
|------------|------------|
| 0.04313859 | 0.02919501 |

Multiple Logistic Regression

```
summary(model3 <- glm(dflt ~ balance + is_student, data = data.default, family = "binomial"))
```

Call:

```
glm(formula = dflt ~ balance + is_student, family = "binomial",  
     data = data.default)
```

Deviance Residuals:

| Min | 1Q | Median | 3Q | Max |
|---------|---------|---------|---------|--------|
| -2.4578 | -0.1422 | -0.0559 | -0.0203 | 3.7435 |

Coefficients:

| | Estimate | Std. Error | z value | Pr(> z) |
|-------------|------------|------------|---------|--------------|
| (Intercept) | -1.075e+01 | 3.692e-01 | -29.116 | < 2e-16 *** |
| balance | 5.738e-03 | 2.318e-04 | 24.750 | < 2e-16 *** |
| is_student | -7.149e-01 | 1.475e-01 | -4.846 | 1.26e-06 *** |

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

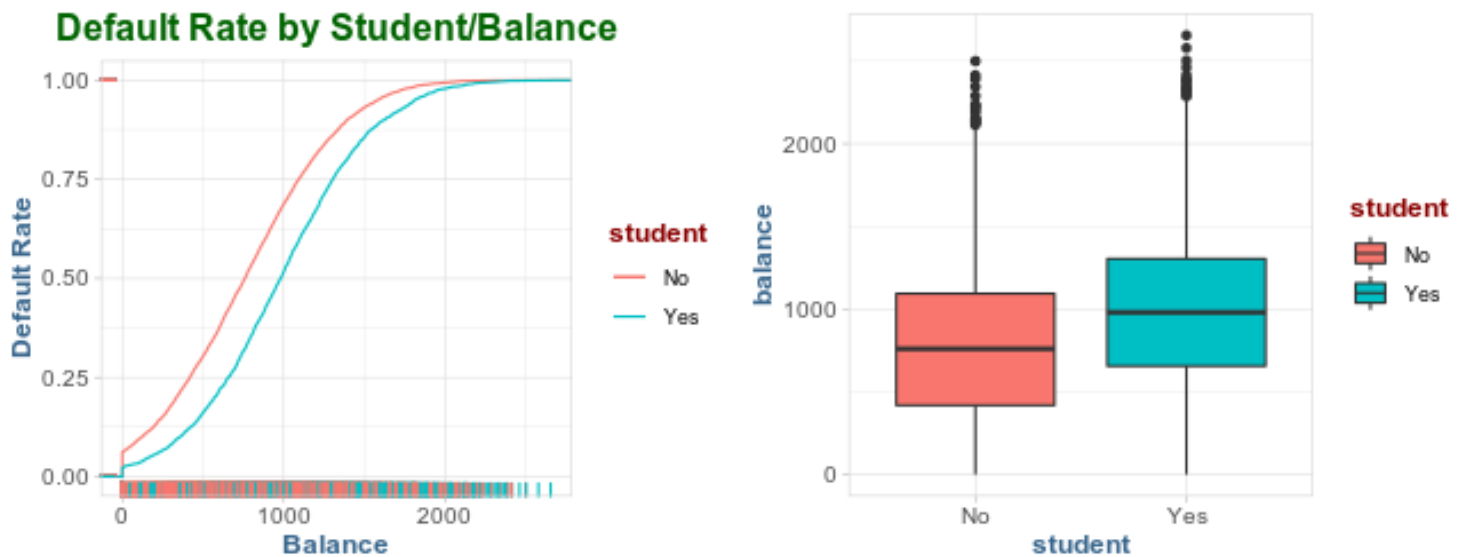
Null deviance: 2920.6 on 9999 degrees of freedom
 Residual deviance: 1571.7 on 9997 degrees of freedom
 AIC: 1577.7

Number of Fisher Scoring iterations: 8

```
p1 <- ggplot(data.default, aes(balance, dflt, color = student)) +
  stat_ecdf() +
  geom_rug(aes(balance, dflt)) +
  labs(x = "Balance", y = "Default Rate", title = "Default Rate by Student/Balance")

p2 <- ggplot(data.default, aes(student, balance, fill = student)) +
  geom_boxplot()

grid.arrange(p1, p2, nrow = 1)
```



```
predict(model3, newdata =
  data.frame( balance = c(1500, 1500),
               is_student = c(1, 0) ),
  type = "response")
```

```
1      2
0.05430945 0.10504923
```

R Lab

```
Smarket <- as.data.table(ISLR::Smarket)
```

```
names(Smarket)
```

```
[1] "Year"      "Lag1"      "Lag2"      "Lag3"      "Lag4"      "Lag5"
[7] "Volume"    "Today"     "Direction"
```

```
dim(Smarket)
```

```
[1] 1250      9
```

```
summary(Smarket)
```

| Year | Lag1 | Lag2 | Lag3 |
|--------------|--------------------|--------------------|--------------------|
| Min. :2001 | Min. :-4.922000 | Min. :-4.922000 | Min. :-4.922000 |
| 1st Qu.:2002 | 1st Qu.: -0.639500 | 1st Qu.: -0.639500 | 1st Qu.: -0.640000 |
| Median :2003 | Median : 0.039000 | Median : 0.039000 | Median : 0.038500 |
| Mean :2003 | Mean : 0.003834 | Mean : 0.003919 | Mean : 0.001716 |
| 3rd Qu.:2004 | 3rd Qu.: 0.596750 | 3rd Qu.: 0.596750 | 3rd Qu.: 0.596750 |
| Max. :2005 | Max. : 5.733000 | Max. : 5.733000 | Max. : 5.733000 |

| Lag4 | Lag5 | Volume | Today |
|--------------------|-------------------|----------------|--------------------|
| Min. :-4.922000 | Min. :-4.92200 | Min. :0.3561 | Min. :-4.922000 |
| 1st Qu.: -0.640000 | 1st Qu.: -0.64000 | 1st Qu.:1.2574 | 1st Qu.: -0.639500 |
| Median : 0.038500 | Median : 0.03850 | Median :1.4229 | Median : 0.038500 |
| Mean : 0.001636 | Mean : 0.00561 | Mean :1.4783 | Mean : 0.003138 |
| 3rd Qu.: 0.596750 | 3rd Qu.: 0.59700 | 3rd Qu.:1.6417 | 3rd Qu.: 0.596750 |
| Max. : 5.733000 | Max. : 5.73300 | Max. :3.1525 | Max. : 5.733000 |

Direction
Down:602
Up :648

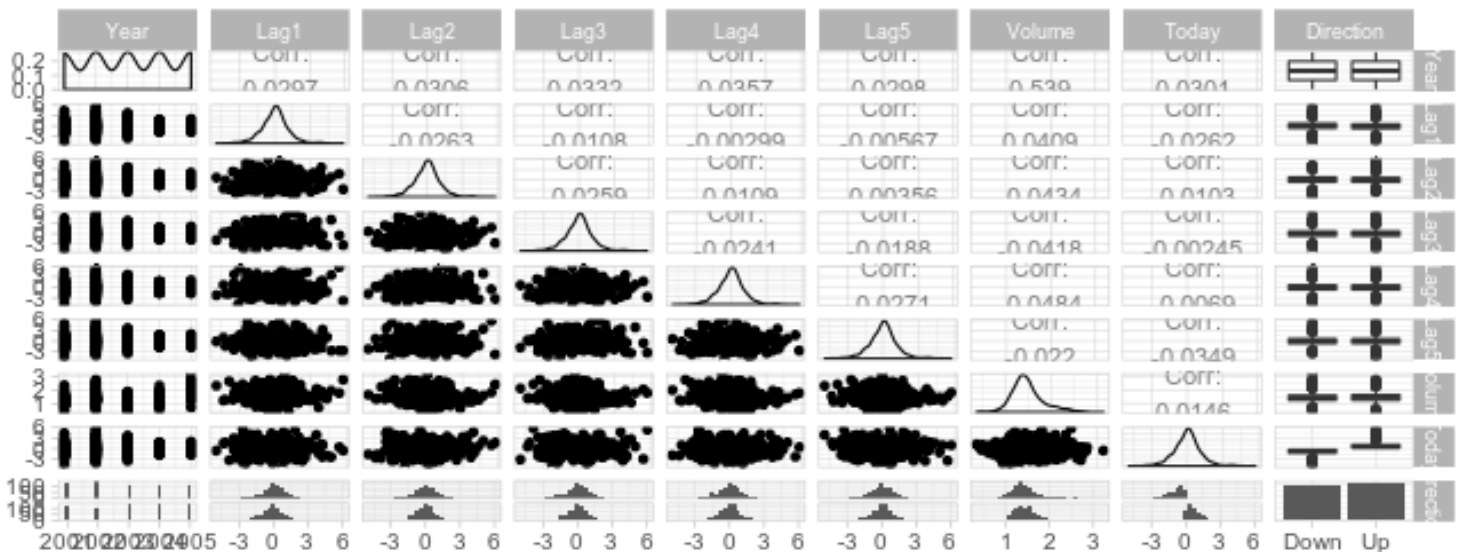
Pairs

```
ggpairs(Smarket) %>%  
  print(progress = F)
```

Warning in ggmatrix_gtable(x, ...): Please use the 'progress' parameter in your ggmatrix-like function call. See ?ggmatrix_progress for a few examples. ggmatrix_gtable 'progress' and 'progress_format' will soon be deprecated.TRUE

```
`stat_bin()` using `bins = 30`. Pick better value with `binwidth`.  
`stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

```
`stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
`stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
`stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
`stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
`stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
`stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
cor(Smarket %>% select(-Direction))
```

| | Year | Lag1 | Lag2 | Lag3 | Lag4 |
|--------|------------|--------------|--------------|--------------|--------------|
| Year | 1.00000000 | 0.029699649 | 0.030596422 | 0.033194581 | 0.035688718 |
| Lag1 | 0.02969965 | 1.000000000 | -0.026294328 | -0.010803402 | -0.002985911 |
| Lag2 | 0.03059642 | -0.026294328 | 1.000000000 | -0.025896670 | -0.010853533 |
| Lag3 | 0.03319458 | -0.010803402 | -0.025896670 | 1.000000000 | -0.024051036 |
| Lag4 | 0.03568872 | -0.002985911 | -0.010853533 | -0.024051036 | 1.000000000 |
| Lag5 | 0.02978799 | -0.005674606 | -0.003557949 | -0.018808338 | -0.027083641 |
| Volume | 0.53900647 | 0.040909908 | -0.043383215 | -0.041823686 | -0.048414246 |
| Today | 0.03009523 | -0.026155045 | -0.010250033 | -0.002447647 | -0.006899527 |

| | Lag5 | Volume | Today |
|--------|--------------|-------------|--------------|
| Year | 0.029787995 | 0.53900647 | 0.030095229 |
| Lag1 | -0.005674606 | 0.04090991 | -0.026155045 |
| Lag2 | -0.003557949 | -0.04338321 | -0.010250033 |
| Lag3 | -0.018808338 | -0.04182369 | -0.002447647 |
| Lag4 | -0.027083641 | -0.04841425 | -0.006899527 |
| Lag5 | 1.000000000 | -0.02200231 | -0.034860083 |
| Volume | -0.022002315 | 1.000000000 | 0.014591823 |
| Today | -0.034860083 | 0.01459182 | 1.000000000 |

```
ggplot(Smarket) +
  geom_boxplot(aes(Year, Volume, group = Year))
```



Logistic Regression

```
summary(glm.fits <- glm(Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 + Volume,
  data = Smarket, family = binomial))
```

Call:

```
glm(formula = Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 +
  Volume, family = binomial, data = Smarket)
```

Deviance Residuals:

| Min | 1Q | Median | 3Q | Max |
|--------|--------|--------|-------|-------|
| -1.446 | -1.203 | 1.065 | 1.145 | 1.326 |

Coefficients:

| | Estimate | Std. Error | z value | Pr(> z) |
|-------------|-----------|------------|---------|----------|
| (Intercept) | -0.126000 | 0.240736 | -0.523 | 0.601 |
| Lag1 | -0.073074 | 0.050167 | -1.457 | 0.145 |
| Lag2 | -0.042301 | 0.050086 | -0.845 | 0.398 |
| Lag3 | 0.011085 | 0.049939 | 0.222 | 0.824 |
| Lag4 | 0.009359 | 0.049974 | 0.187 | 0.851 |

| | | | | |
|--------|----------|----------|-------|-------|
| Lag5 | 0.010313 | 0.049511 | 0.208 | 0.835 |
| Volume | 0.135441 | 0.158360 | 0.855 | 0.392 |

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1731.2 on 1249 degrees of freedom
 Residual deviance: 1727.6 on 1243 degrees of freedom
 AIC: 1741.6

Number of Fisher Scoring iterations: 3

```
coef(glm.fits)
```

| (Intercept) | Lag1 | Lag2 | Lag3 | Lag4 | Lag5 |
|--------------|--------------|--------------|-------------|-------------|-------------|
| -0.126000257 | -0.073073746 | -0.042301344 | 0.011085108 | 0.009358938 | 0.010313068 |
| Volume | | | | | |
| 0.135440659 | | | | | |

Probabilites of going up (first 10 trading days)

```
glm.probs <- predict(glm.fits, type = "response")
head(glm.probs, 10)
```

| 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 |
|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|
| 0.5070841 | 0.4814679 | 0.4811388 | 0.5152224 | 0.5107812 | 0.5069565 | 0.4926509 | 0.5092292 |
| 9 | 10 | | | | | | |
| 0.5176135 | 0.4888378 | | | | | | |

```
contrasts(Smarket$Direction)
```

| | |
|------|----|
| | Up |
| Down | 0 |
| Up | 1 |

Predictions

```
glm.pred <- rep("Down", nrow(Smarket))
glm.pred[glm.probs > 0.5] <- "Up"
```

```
table(glm.pred, Smarket$Direction)
```

| | | |
|----------|------|-----|
| glm.pred | Down | Up |
| Down | 145 | 141 |
| Up | 457 | 507 |

```
mean(glm.pred == Smarket$Direction)
```

```
[1] 0.5216
```

Validation

Get the holdout set.

```
train <- (Smarket$Year < 2005)
Smarket.2005 <- Smarket[!train]
dim(Smarket.2005)
```

```
[1] 252  9
```

```
Direction.2005 <- Smarket$Direction[!train]
```

Train the logistic regression model.

```
glm.fits <- glm(Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 + Volume,
                 data = Smarket, family = binomial, subset = train)

glm.probs <- predict(glm.fits, Smarket.2005, type = "response")
```

Test

```
glm.pred <- rep("Down", 252)
glm.pred[glm.probs > 0.5] <- "Up"

table(glm.pred, Direction.2005)
```

```
      Direction.2005
glm.pred Down Up
Down    77  97
Up      34  44
```

```
mean(glm.pred == Direction.2005)
```

```
[1] 0.4801587
```

Model 2

```
summary(glm.fits <- glm(Direction ~ Lag1 + Lag2, data = Smarket, family = binomial, subset = train))
```

Call:

```
glm(formula = Direction ~ Lag1 + Lag2, family = binomial, data = Smarket,
     subset = train)
```

Deviance Residuals:

```
      Min       1Q   Median       3Q      Max
-1.345  -1.188   1.074   1.164   1.326
```

Coefficients:

```
      Estimate Std. Error z value Pr(>|z|)
```



```
(Intercept)  0.03222    0.06338    0.508    0.611
Lag1         -0.05562    0.05171   -1.076    0.282
Lag2         -0.04449    0.05166   -0.861    0.389
```

(Dispersion parameter for binomial family taken to be 1)

```
Null deviance: 1383.3  on 997  degrees of freedom
Residual deviance: 1381.4  on 995  degrees of freedom
AIC: 1387.4
```

Number of Fisher Scoring iterations: 3

```
glm.probs <- predict(glm.fits, Smarket.2005, type = "response")
glm.pred <- rep("Down", nrow(Smarket.2005))
```

```
glm.pred[glm.probs >= 0.5] <- "Up"
```

```
table(glm.pred, Direction.2005)
```

```
      Direction.2005
glm.pred Down  Up
Down     35   35
Up       76  106
```

```
mean(glm.pred == Direction.2005)
```

```
[1] 0.5595238
```

```
predict(glm.fits, newdata = data.table(Lag1 = c(1.2, 1.5),
                                         Lag2 = c(1.1, -0.8)),
        type = "response")
```

```
      1      2
0.4791462 0.4960939
```

Linear Discriminant Analysis

LDA is from **MASS** package.

```
summary(lda.fit <- lda(Direction ~ Lag1 + Lag2, data = Smarket, subset = train))
```

```
      Length Class  Mode
prior     2      -none- numeric
counts    2      -none- numeric
means     4      -none- numeric
scaling   2      -none- numeric
lev       2      -none- character
```

```
svd      1      -none- numeric
N        1      -none- numeric
call     4      -none- call
terms    3      terms  call
xlevels  0      -none- list
```

```
lda.fit
```

Call:

```
lda(Direction ~ Lag1 + Lag2, data = Smarket, subset = train)
```

Prior probabilities of groups:

```
      Down      Up
0.491984 0.508016
```

Group means:

```
      Lag1      Lag2
Down 0.04279022 0.03389409
Up   -0.03954635 -0.03132544
```

Coefficients of linear discriminants:

```
      LD1
Lag1 -0.6420190
Lag2 -0.5135293
```

```
lda.pred <- predict(lda.fit, Smarket.2005)
```

```
names(lda.pred)
```

```
[1] "class"      "posterior" "x"
```

Predictions:

```
lda.class <- lda.pred$class
table(lda.class, Direction.2005)
```

```
      Direction.2005
lda.class Down  Up
      Down   35  35
      Up    76 106
```

Note: almost identical to logistic regression.

```
mean(lda.class == Direction.2005)
```

```
[1] 0.5595238
```

```
sum(lda.pred$posterior[, 1] >= 0.5)
```

```
[1] 70
```

```
sum(lda.pred$posterior[, 1] < 0.5)
```

```
[1] 182
```

The posterior probabilities output by the model corresponds to the probability that the market will decrease.

```
lda.pred$posterior[1:20, 1]
```

| 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 |
|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|
| 0.4901792 | 0.4792185 | 0.4668185 | 0.4740011 | 0.4927877 | 0.4938562 | 0.4951016 | 0.4872861 |
| 9 | 10 | 11 | 12 | 13 | 14 | 15 | 16 |
| 0.4907013 | 0.4844026 | 0.4906963 | 0.5119988 | 0.4895152 | 0.4706761 | 0.4744593 | 0.4799583 |
| 17 | 18 | 19 | 20 | | | | |
| 0.4935775 | 0.5030894 | 0.4978806 | 0.4886331 | | | | |

```
lda.class[1:20]
```

```
[1] Up Up Up Up Up Up Up Up Up Up Up Down Up Up Up
[16] Up Up Down Up Up
Levels: Down Up
```

Apply a threshold of 90% to predictions:

```
sum(lda.pred$posterior[, 1] > .9)
```

```
[1] 0
```

Quadratic Discriminant Analysis

```
summary(qda.fit <- qda(Direction ~ Lag1 + Lag2, data = Smarket, subset = train))
```

| | Length | Class | Mode |
|---------|--------|--------|-----------|
| prior | 2 | -none- | numeric |
| counts | 2 | -none- | numeric |
| means | 4 | -none- | numeric |
| scaling | 8 | -none- | numeric |
| ldet | 2 | -none- | numeric |
| lev | 2 | -none- | character |
| N | 1 | -none- | numeric |
| call | 4 | -none- | call |
| terms | 3 | terms | call |
| xlevels | 0 | -none- | list |

```
qda.fit
```

Call:

```
qda(Direction ~ Lag1 + Lag2, data = Smarket, subset = train)
```

Prior probabilities of groups:

| | Down | Up |
|--|----------|----------|
| | 0.491984 | 0.508016 |

Group means:

| | Lag1 | Lag2 |
|------|-------------|-------------|
| Down | 0.04279022 | 0.03389409 |
| Up | -0.03954635 | -0.03132544 |

Predictions

```
qda.class <- predict(qda.fit, Smarket.2005)$class
table(qda.class, Direction.2005)
```

| | Direction.2005 | |
|-----------|----------------|-----|
| qda.class | Down | Up |
| Down | 30 | 20 |
| Up | 81 | 121 |

```
mean(qda.class == Direction.2005)
```

```
[1] 0.5992063
```

K-Nearest Neighbors

Data Setup

```
train.X <- with(Smarket, cbind(Lag1, Lag2))[train, ]
test.X <- with(Smarket, cbind(Lag1, Lag2))[!train, ]
train.Direction <- Smarket$Direction[train]
```

KNN

```
set.seed(1)
```

```
knn.pred <- knn(train.X, test.X, train.Direction, k = 1)
table(knn.pred, Direction.2005)
```

| | Direction.2005 | |
|----------|----------------|----|
| knn.pred | Down | Up |
| Down | 43 | 58 |
| Up | 68 | 83 |

```
set.seed(1)
```

```
knn.pred <- knn(train.X, test.X, train.Direction, k = 3)
table(knn.pred, Direction.2005)
```

```

      Direction.2005
knn.pred Down Up
      Down   48 55
      Up     63 86

```

Caravan Insurance Data

```
caravan <- Caravan
```

```
dim(caravan)
```

```
[1] 5822  86
```

```
summary(caravan$Purchase)
```

```

      No   Yes
5474   348

```

```
table(caravan$Purchase) %>% prop.table()
```

```

      No      Yes
0.94022673 0.05977327

```

```
standardized.X <- scale(caravan[, -86])
```

```
var(caravan[, 1])
```

```
[1] 165.0378
```

```
var(caravan[, 2])
```

```
[1] 0.1647078
```

```
var(standardized.X[, 1])
```

```
[1] 1
```

```
var(standardized.X[, 2])
```

```
[1] 1
```

- K=1

```
test <- 1:1000
```

```
train.X <- standardized.X[-test,]
```

```
test.X <- standardized.X[test,]
```

```
train.Y <- caravan$Purchase[-test]
```

```
test.Y <- caravan$Purchase[test]
```

```
set.seed(1)

knn.pred <- knn(train.X, test.X, train.Y, k = 1)
mean(test.Y != knn.pred)
```

```
[1] 0.118
```

```
mean(test.Y != "No")
```

```
[1] 0.059
```

```
result <- table(knn.pred, test.Y)
result
```

```
      test.Y
knn.pred No Yes
      No  873  50
      Yes  68   9
```

```
result %>% prop.table()
```

```
      test.Y
knn.pred   No   Yes
      No  0.873 0.050
      Yes 0.068 0.009
```

- K=3

```
set.seed(1)

knn.pred <- knn(train.X, test.X, train.Y, k = 3)
mean(test.Y != knn.pred)
```

```
[1] 0.074
```

```
mean(test.Y != "No")
```

```
[1] 0.059
```

```
result <- table(knn.pred, test.Y)
result
```

```
      test.Y
knn.pred No Yes
      No  921  54
      Yes  20   5
```

```
result %>% prop.table()
```

```
      test.Y
```

```
knn.pred    No   Yes
      No 0.921 0.054
      Yes 0.020 0.005
```

- K=5

```
set.seed(1)
```

```
knn.pred <- knn(train.X, test.X, train.Y, k = 5)
mean(test.Y != knn.pred)
```

```
[1] 0.066
```

```
mean(test.Y != "No")
```

```
[1] 0.059
```

```
result <- table(knn.pred, test.Y)
result
```

```
      test.Y
knn.pred No Yes
      No  930  55
      Yes   11   4
```

```
result %>% prop.table()
```

```
      test.Y
knn.pred    No   Yes
      No 0.930 0.055
      Yes 0.011 0.004
```

Logistic Regression Alternative

```
glm.fits <- glm(Purchase ~ ., data = caravan, family = binomial,
               subset = -test)
```

Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

```
glm.probs <- predict(glm.fits, caravan[test,], type = "response")
```

```
# .5 cut-off
```

```
glm.pred <- rep("No", 1000)
glm.pred[glm.probs > .5] <- "Yes"
```

```
results <- table(glm.pred, test.Y)
results %>% prop.table()
```

```
      test.Y
glm.pred    No   Yes
      No 0.934 0.059
```

```
Yes 0.007 0.000
```

```
# .25 cut-off
glm.pred <- rep("No", 1000)
glm.pred[glm.probs > .25] <- "Yes"

results <- table(glm.pred, test.Y)
results %>% prop.table()
```

```
      test.Y
glm.pred No  Yes
No      0.919 0.048
Yes     0.022 0.011
```

```
# Quiz

bal <- 1936.75

exp(-10.6513 + 0.0055 * bal) / ( 1 + exp(-10.6513 + 0.0055*bal))
```

```
[1] 0.5002062

b0 <- -6; b1 <- 0.05; b2 <- 1
x1 <- 50; x2 <- 3.5
exp(b0 + b1 * x1 + b2 * x2) / ( 1 + exp(b0 + b1 * x1 + b2 * x2))

[1] 0.5
```

Conceptual

- 1.) Using a little bit of algebra, prove that (4.2) is equivalent to (4.3). In other words, the logistic function representation and logit representation for the logistic regression models are equivalent
- 2.) It was stated in the text that classifying an observation to the class for which (4.13) is largest. Prove that this is the case. In other words, under the assumption that the observations in the k th class are drawn from a $N(\mu_k, \sigma_k^2)$ distribution, the Bayes' classifier assigns an observation to the class for which the discriminant function is maximized.
- 3.) This problem relates to the QDA model, in which the observations within each class are drawn from a normal distribution with a class-specific mean vector and a class specific covariance matrix. We consider the simple case where $p=1$; i.e. there is only one feature.

Suppose that we have K classes, and if an observation belongs to the k th class then X comes from a one-dimensional normal distribution, $X \sim N(\mu_k, \sigma_k^2)$. Recall that the density function for the one-dimensional normal distribution is given in (4.11). Prove that in this case, the Bayes' classifier is not linear. Argue that it is in fact quadratic.

Applied**8.)**